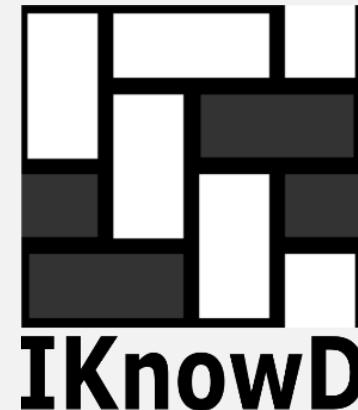


# MADEIRA INTERNATIONAL WORKSHOP IN MACHINE LEARNING

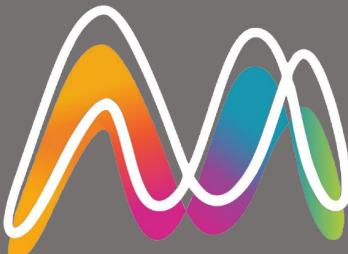


Google Developer Groups  
Madeira



## ORDEM DOS ENGENHEIROS TÉCNICOS

Platinum sponsor:



malogica solutions

Silver sponsor:



Institutional support

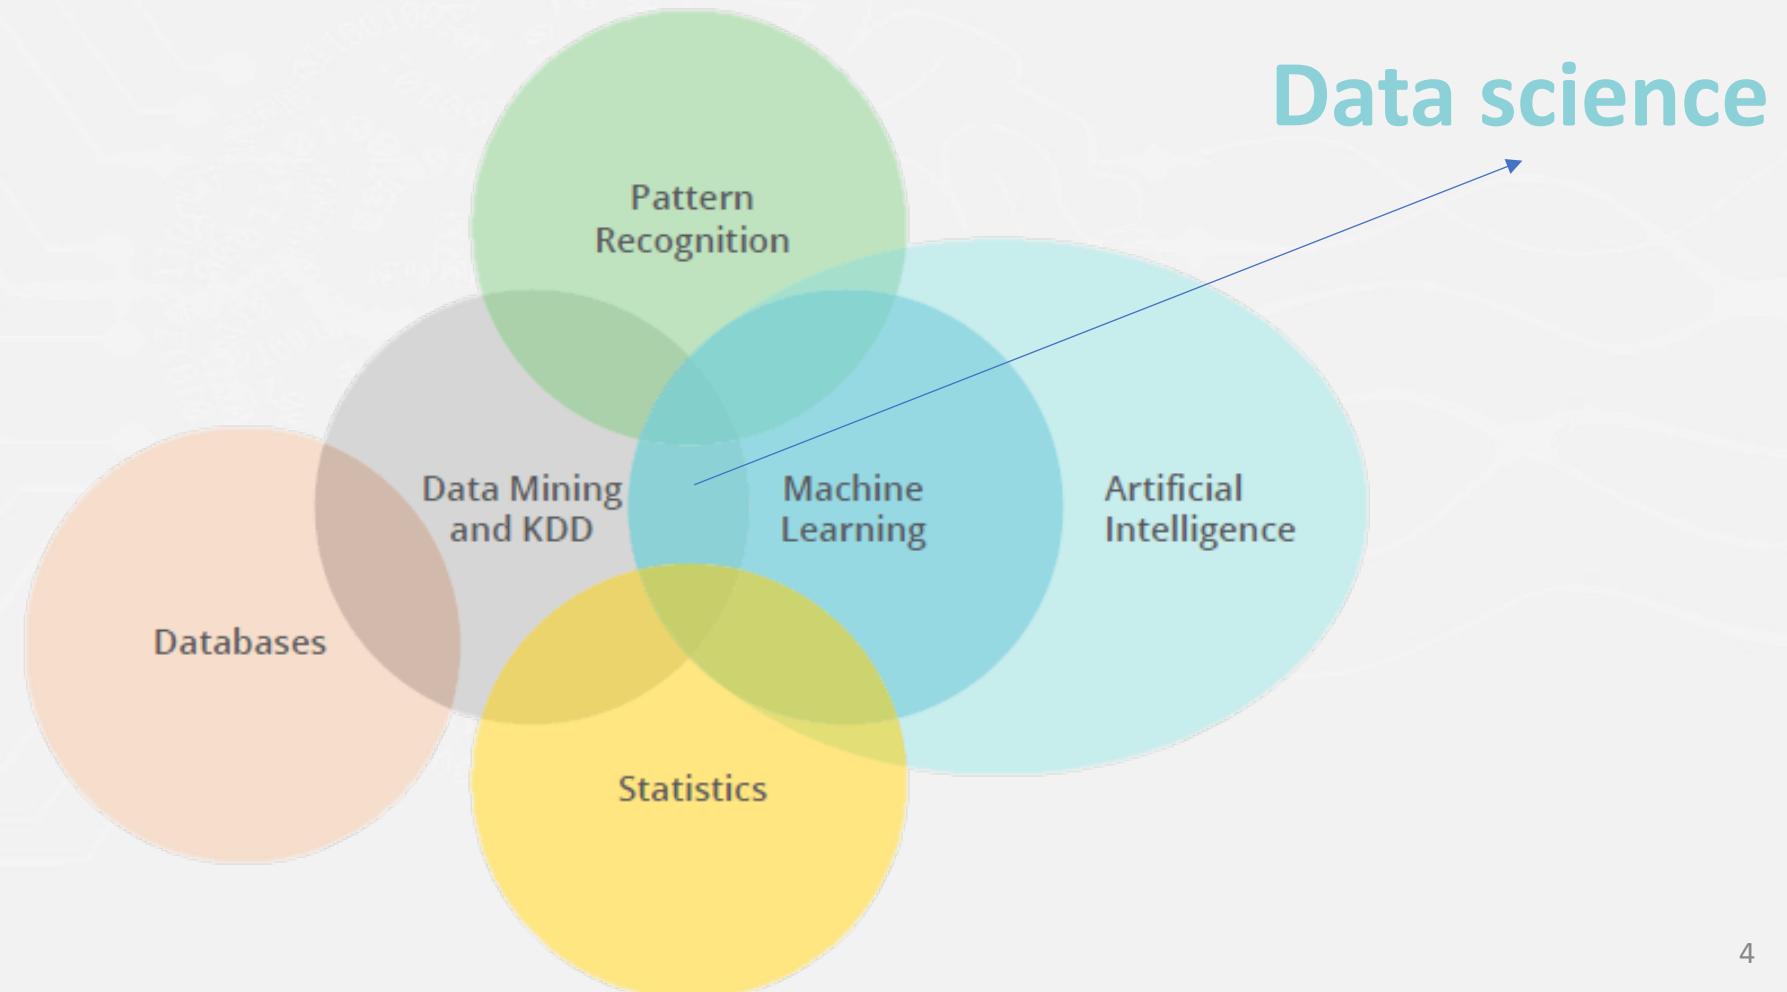


# Outline

- 1- What is machine learning?
- 2- Examples of machine learning applications
- 3- Key remarks

# 1- What is machine learning?

# 1- What is machine learning?



# 1- What is machine learning?

“Field of study that **gives computers the ability to learn** without being explicitly programmed”

Arthur Samuel, 1959

“A **computer program** is said to **learn from experience E** with respect to some task **T** and some performance measure **P** if its performance on **T**, as measured by **P**, **improves with experience E**”

Tom Mitchell, 1997

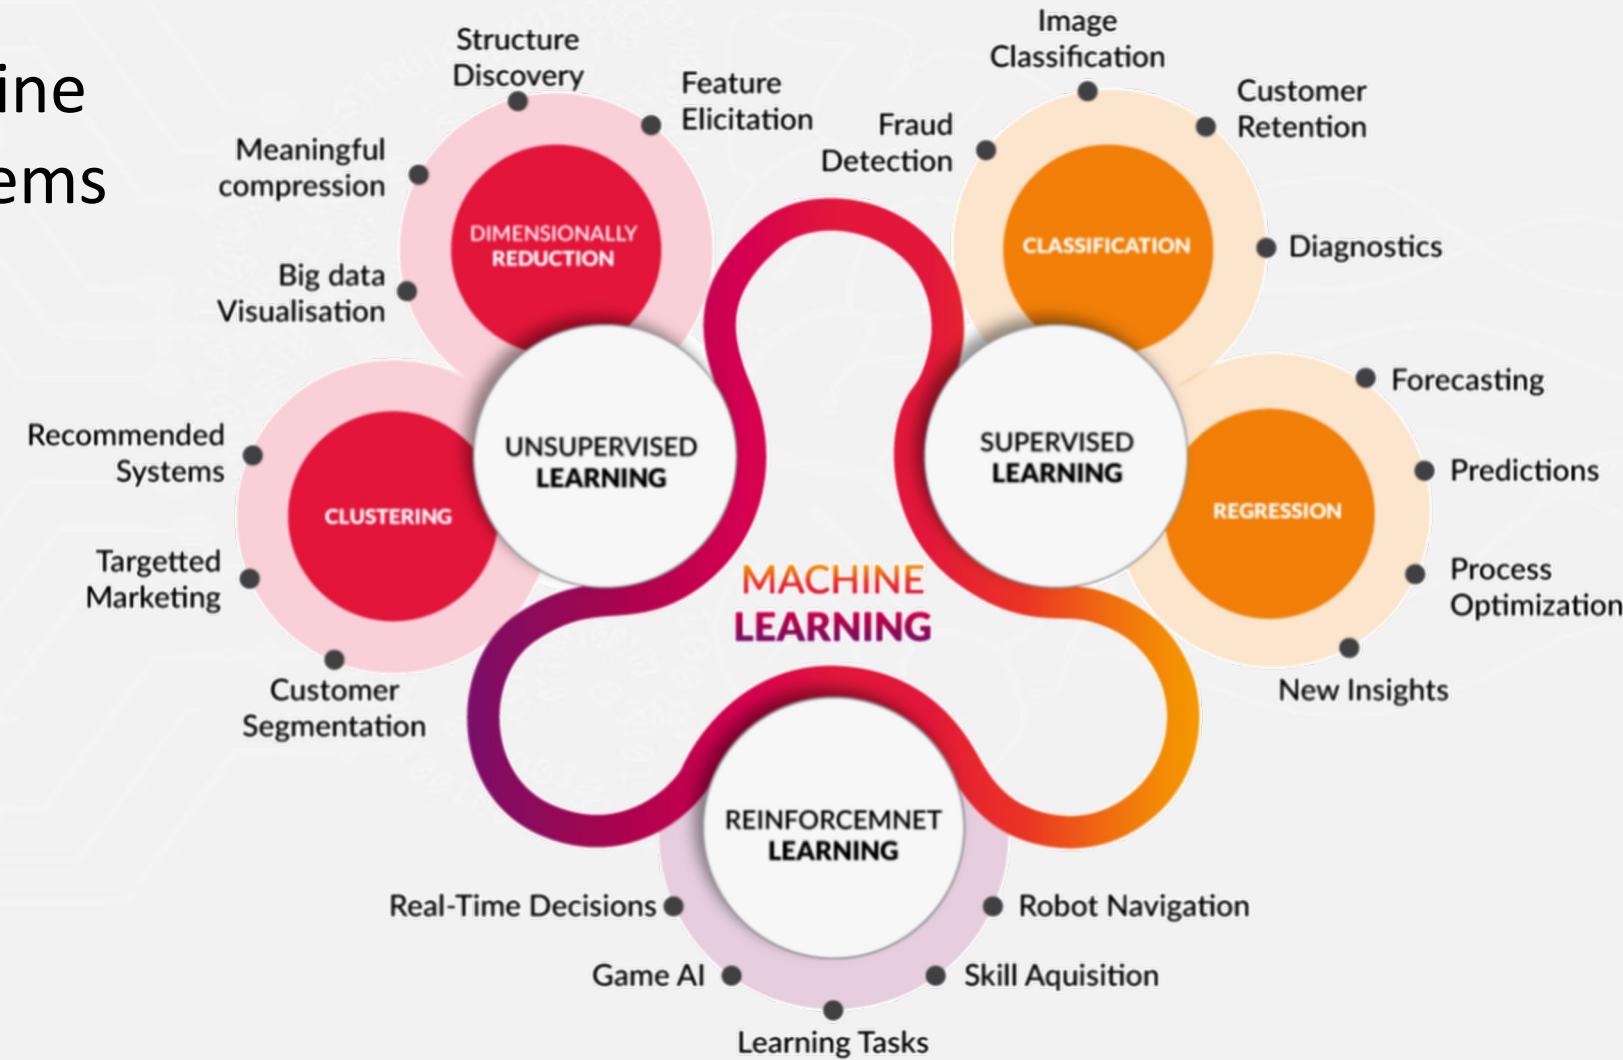
Using data for answering questions

Training

Predicting

# 1- What is machine learning?

Types of machine learning problems



# 1- What is machine learning?

Supervised

Learn through **examples** of which we know the desired output (what we want to predict).

Unsupervised

-Is this an image of a cat, a dog or a horse?

-Are these emails spam or not?

Reinforcement

-What will be the energy production of this windfarm?

# 1- What is machine learning?

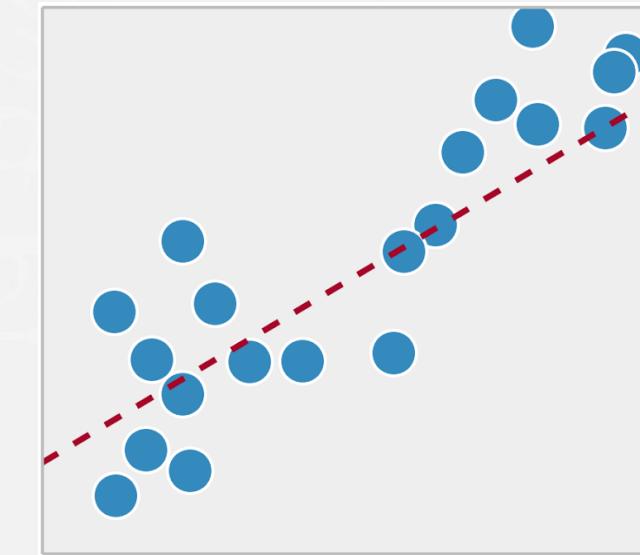
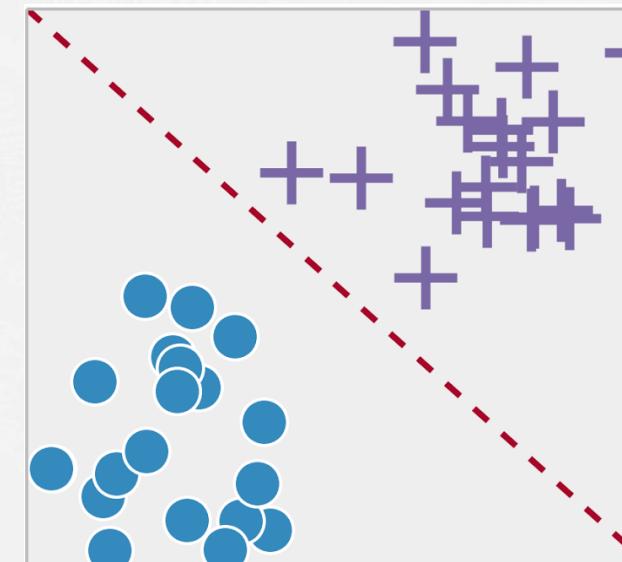
Supervised

Classification

Regression

Unsupervised

Reinforcement



# 1- What is machine learning?

Supervised

Unsupervised

Reinforcement

There is **no desired output** thus it learns something about the data.

-I have set of sports balls and want to divide them into five groups according to how similar they are.

-Learn a representation from a group of photos from a cat.

# 1- What is machine learning?

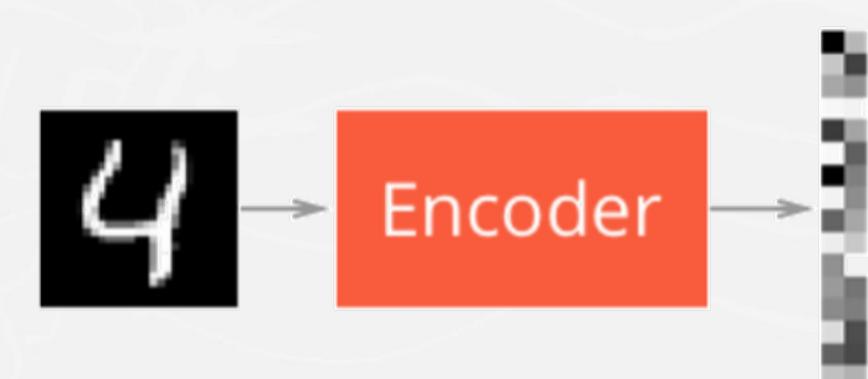
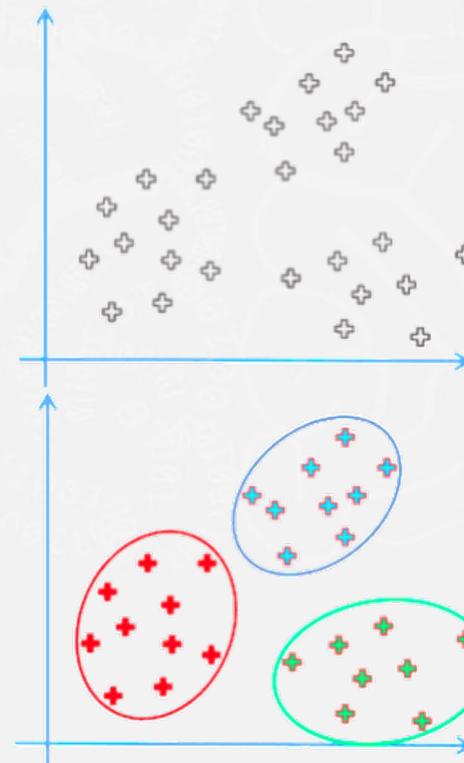
Supervised

Clustering

Dimensionality  
reduction

Unsupervised

Reinforcement



# 1- What is machine learning?

Supervised

Unsupervised

Reinforcement

An agent **interacts** with an **environment** and watches the results of the interaction. The **environment** gives **feedback** via positive or negative **reward signal**.

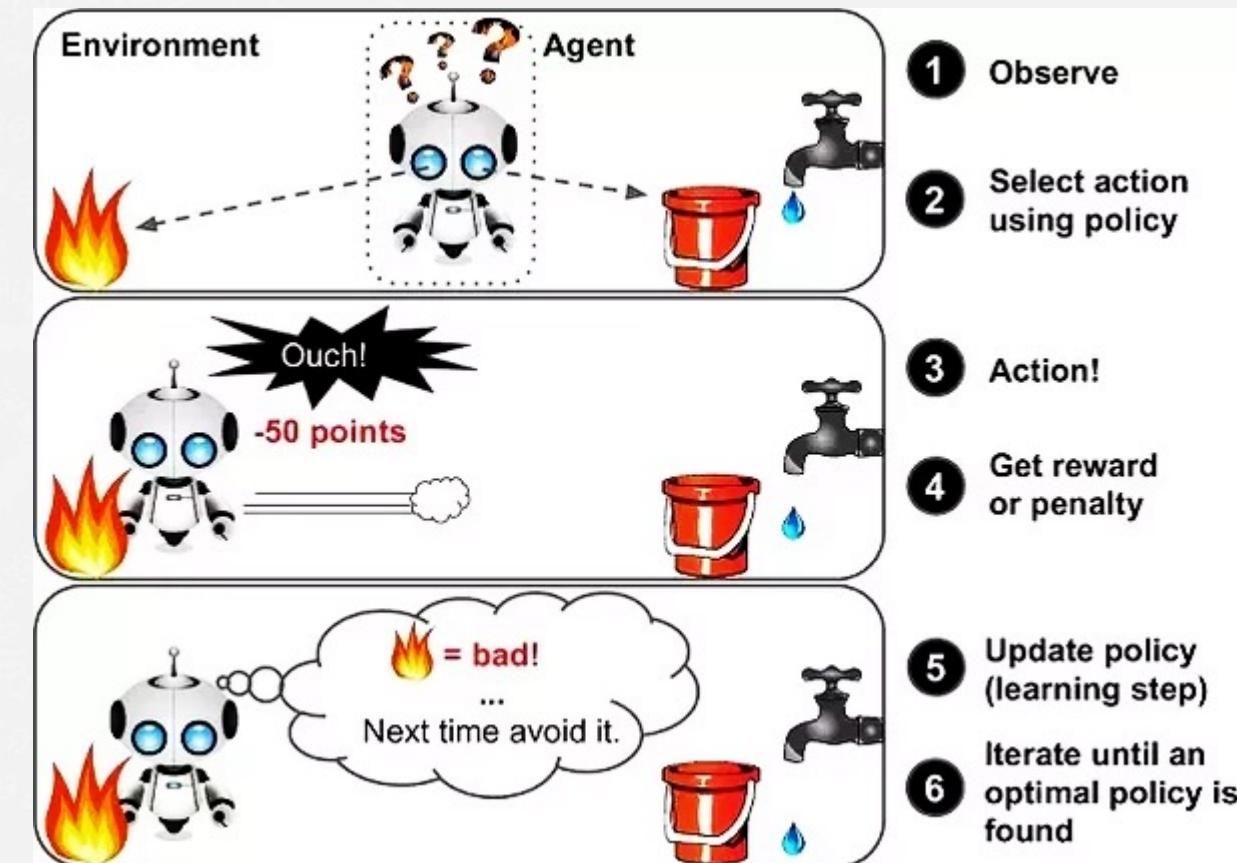
-Decide the next action a character should do in a game.

# 1- What is machine learning?

Supervised

Unsupervised

Reinforcement

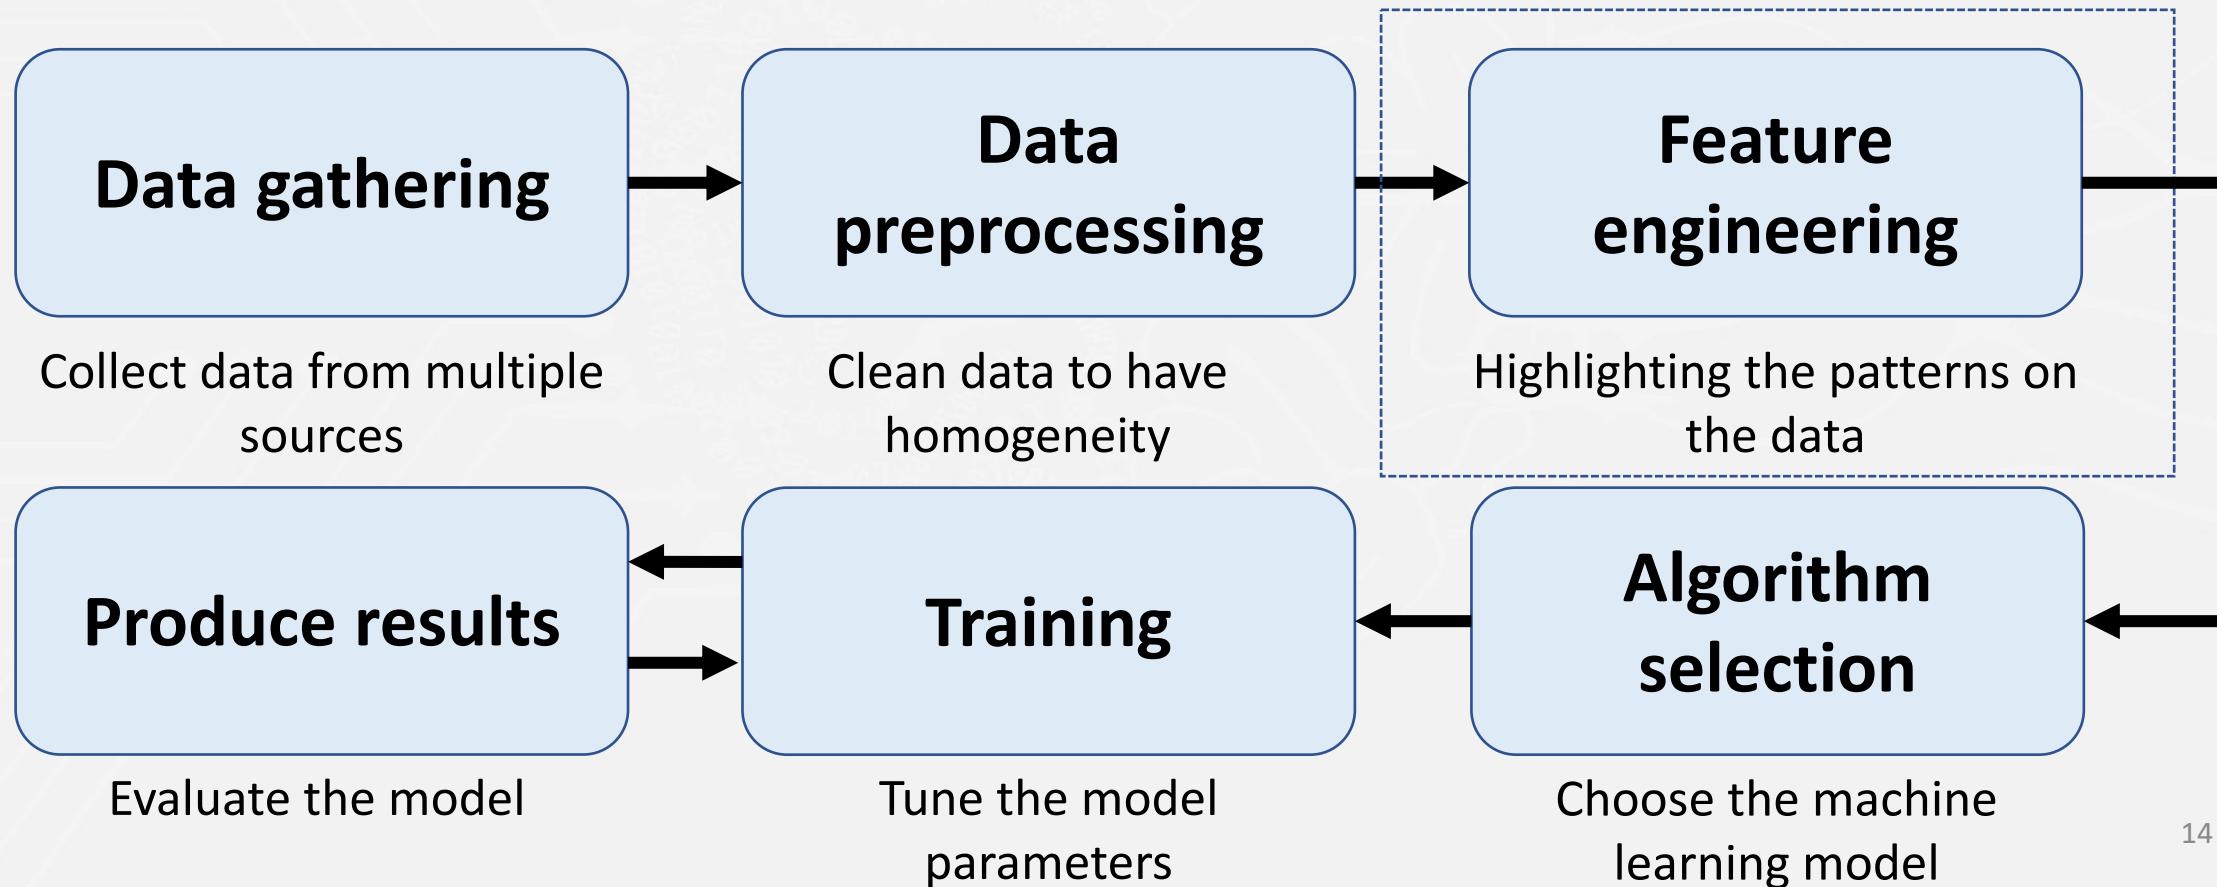


## 2- Examples of machine learning applications

## 2- Examples of machine learning applications

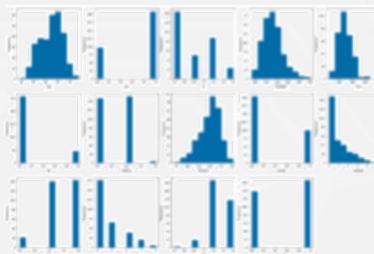
Steps to solve a machine learning problem

Step not mandatory when using deep learning



## Exploratory Data Analysis (EDA)

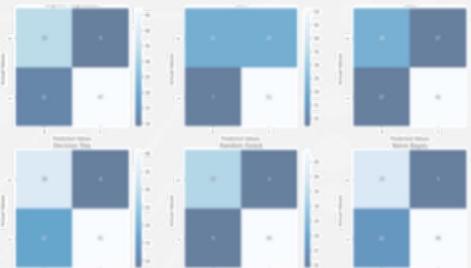
- 1) Histogram: `df.plot(kind = 'hist')`
- 2) Box Plot: `sns.boxplot()`
- 3) Grouped Bar Chart: `sns.countplot()`



## Model Evaluation

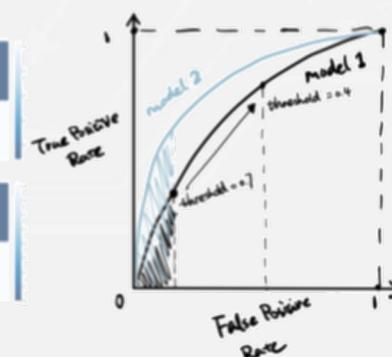
### Confusion Matrix

```
confusion_matrix(y_test, y_pred)
```

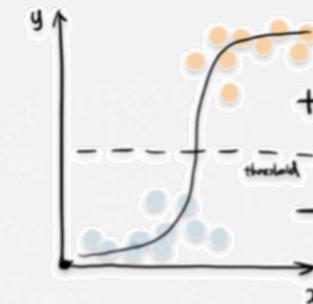


### ROC & AUC

```
metrics.auc(fpr, tpr)
```



## Logistic Regression



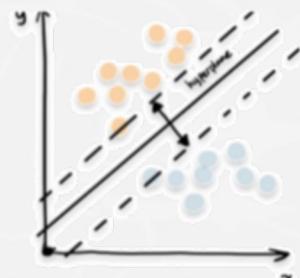
## Decision Tree



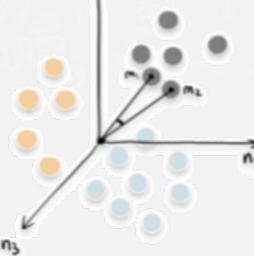
## Random Forest



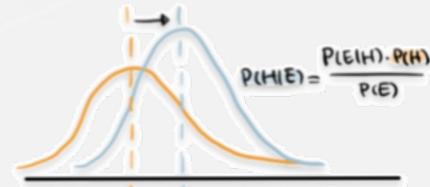
## Support Vector Machine



## K Nearest Neighbour



## Naive Bayes



## Dataset

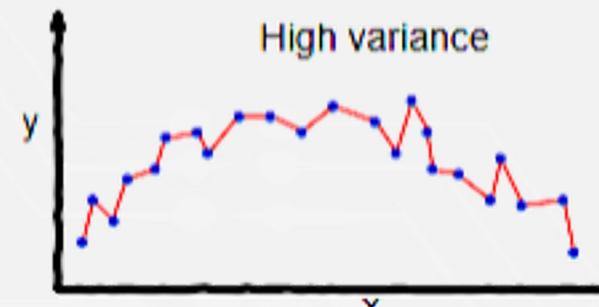
## Training

## Testing

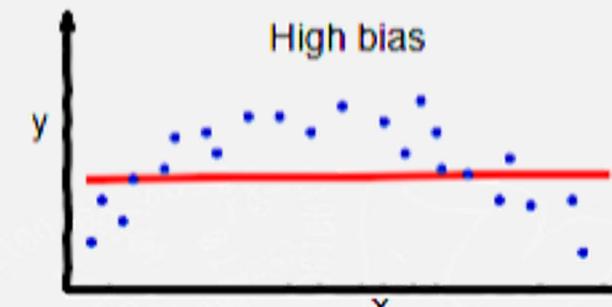
## Training

## Validation

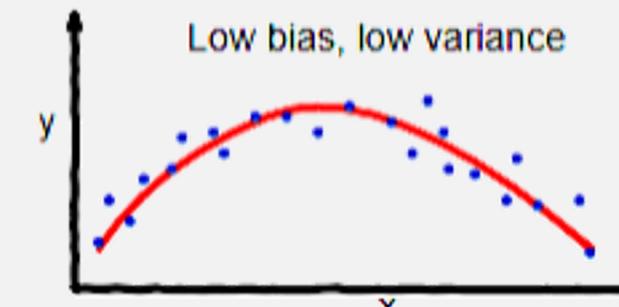
## Testing



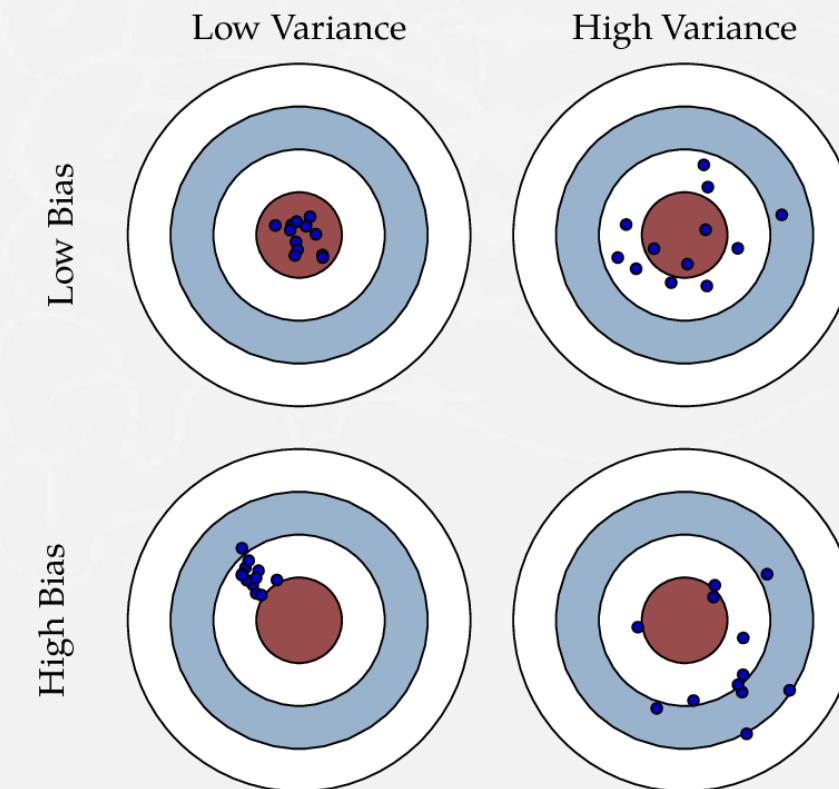
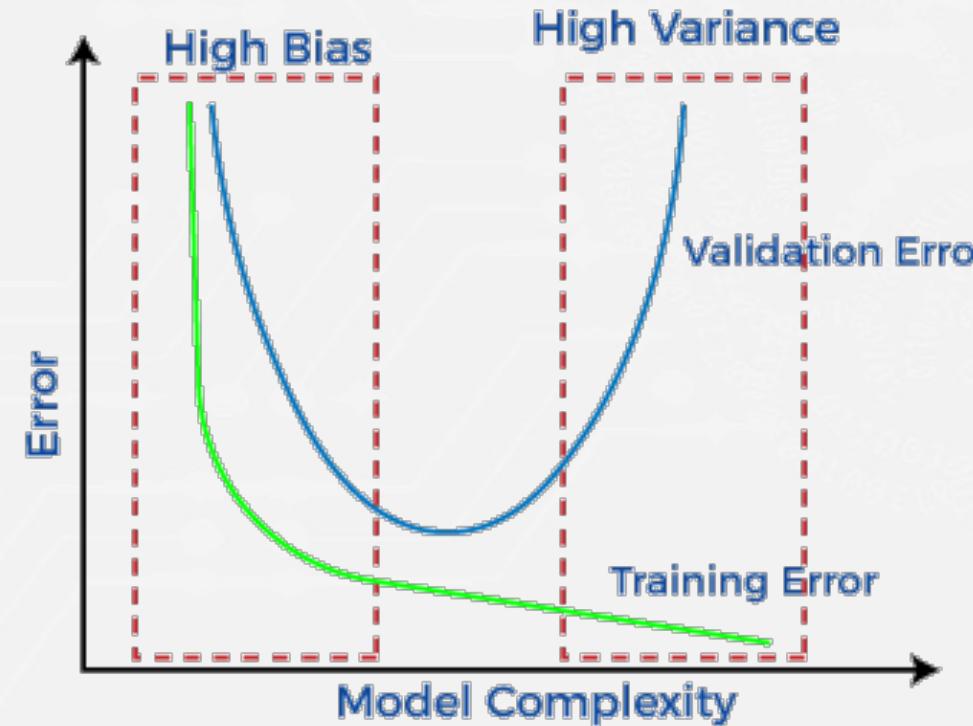
overfitting



underfitting



Good balance



# 2- Examples of machine learning applications

Classification

# 2- Examples of machine learning applications

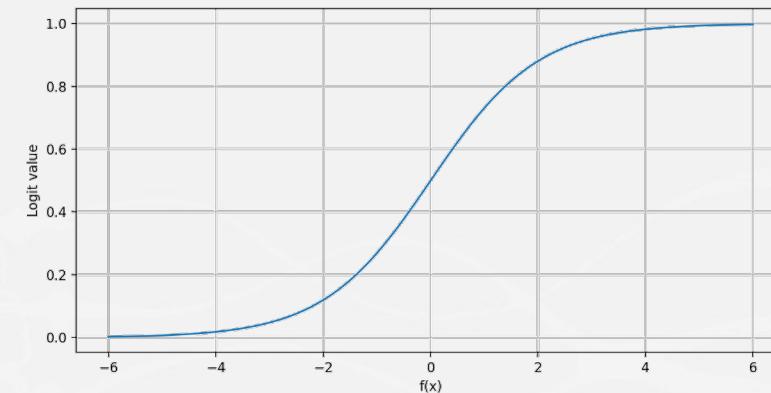
## Classification - Logistic regression

Logistic function:

$$y = \frac{e^{f(x_1, x_2, \dots, x_n)}}{1 + e^{f(x_1, x_2, \dots, x_n)}} = \frac{1}{1 + e^{-f(x_1, x_2, \dots, x_n)}}$$

Where

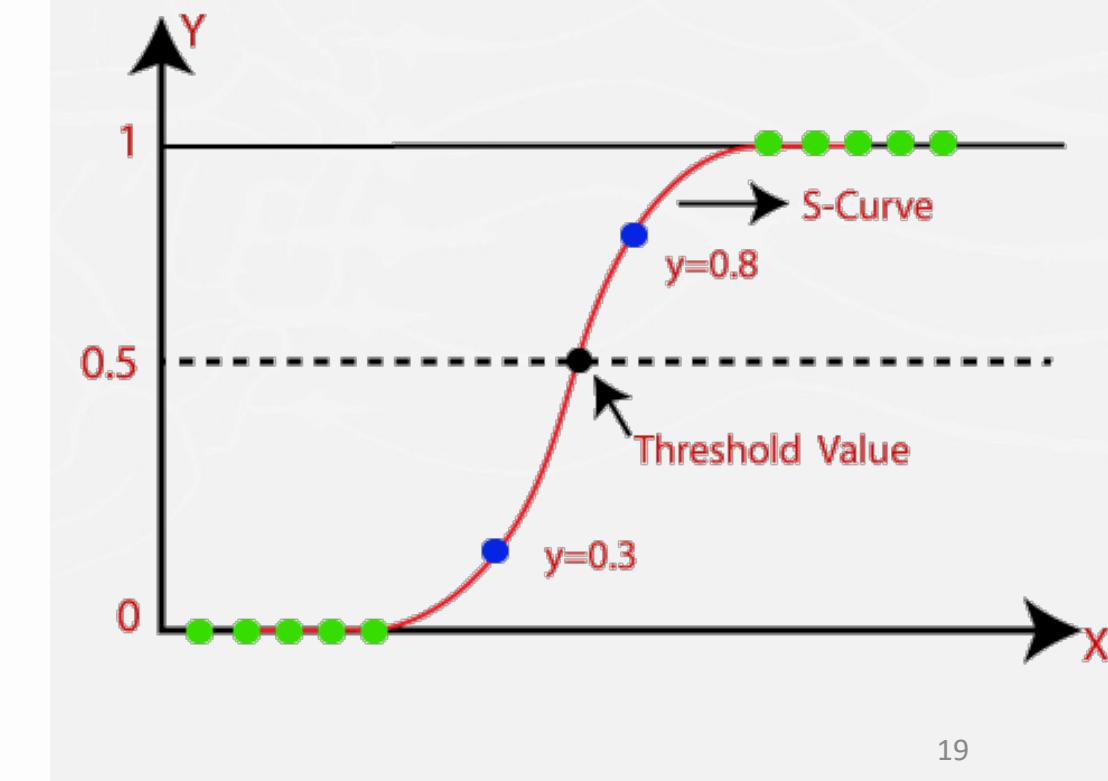
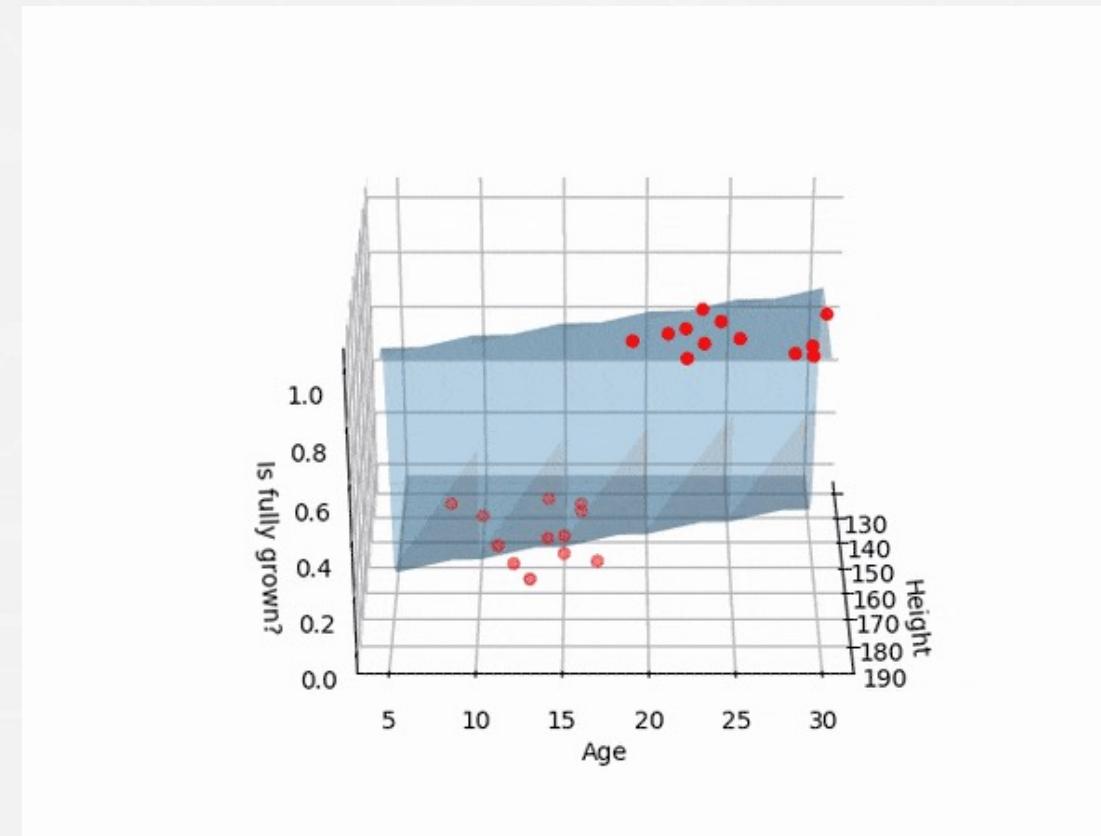
$$f(x_1, x_2, \dots, x_n) = a_0 + a_1 x_1 + \dots + a_n x_n$$



```
1      """
2      Plot Logit function
3      """
4      #Sigmoid Function
5      import matplotlib.pyplot as plt
6      import numpy as np
7      import math
8      def logit(fx):
9          return math.exp(fx)/(1+math.exp(fx))
10     fx = np.linspace(-6,6,100)
11     y = np.vectorize(logit)
12     plt.plot(fx,y(fx))
13     plt.xlabel('f(x)')
14     plt.ylabel('Logit value')
15     plt.grid()
```

## 2- Examples of machine learning applications

### Classification - Logistic regression



## 2- Examples of machine learning applications

### Diabetes prediction

### Pima Indian Diabetes dataset

ATTRIBUTE	DESCRIPTION
Preg	Number of pregnancies
Plas	Plasma glucose concentration in an oral glucose tolerance test
Pres	Diastolic blood pressure
Skin	Triceps skin fold thickness
Insu	2-Hour serum insulin
Mass	Body mass index
Pedi	Diabetes pedigree function
Age	Age of an individual
class	Tested positive / negative

	Mean	Standard deviation	Min/max
Preg	3.8	3.4	1/17
Plas	120.9	32	56/197
Pres	69.1	19.4	24/110
Skin	20.5	16	7/52
Insu	79.8	115.2	15/846
Mass	32	7.9	18.2/57.3
Pedi	0.5	0.3	0.0850/2.3290
Age	33.2	11.8	21/81

# 2- Examples of machine learning applications

## Classification - Logistic regression

```
1  """
2  Loading Data
3  """
4
5  #import pandas
6  import pandas as pd
7  col_names = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
8  'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome']
9
10 # load dataset
11 pima = pd.read_csv("diabetes.csv", header=None, names=col_names)
12 pima = pima.iloc[1:]
13 """
14 Selecting Feature
15 """
16 #split dataset in desired features to be examined and target variable
17 feature_cols = ['Pregnancies', 'Insulin', 'BMI', 'Age','Glucose',
18 'BloodPressure','DiabetesPedigreeFunction']
19 X = pima[feature_cols] # Features
20 y = pima.Outcome # Target variable
21 """
22 Splitting Data
23 """
24 # split X and y into training and testing sets
25 from sklearn.model_selection import train_test_split
26 X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25,
27 random_state=0)
28 """
29 Model Development and Prediction
30 """
31 # import the class
32 from sklearn.linear_model import LogisticRegression
33 # instantiate the model (using default parameters with maximum of 1000 runs)
34 logreg = LogisticRegression(max_iter=1000)
35 # fit the model with data
36 logreg.fit(X_train,y_train)
37 # make predictions on the testing set
38 y_pred=logreg.predict(X_test)
39 """
40 Model Evaluation
41 """
42 from sklearn import metrics
43 y_test = [ int(s) for s in y_test]
44 y_pred = [ int(s) for s in y_pred]
45 cnf_matrix = metrics.confusion_matrix(y_test, y_pred)
46 print(cnf_matrix)
47 print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
48 print("Precision:",metrics.precision_score(y_test, y_pred))
49 print("Recall:",metrics.recall_score(y_test, y_pred))
50 import matplotlib.pyplot as plt
51 y_pred_proba = logreg.predict_proba(X_test)[:,1]
52 fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
53 auc = metrics.roc_auc_score(y_test, y_pred_proba)
54 plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
55 plt.legend(loc=4)
56 plt.show()
```

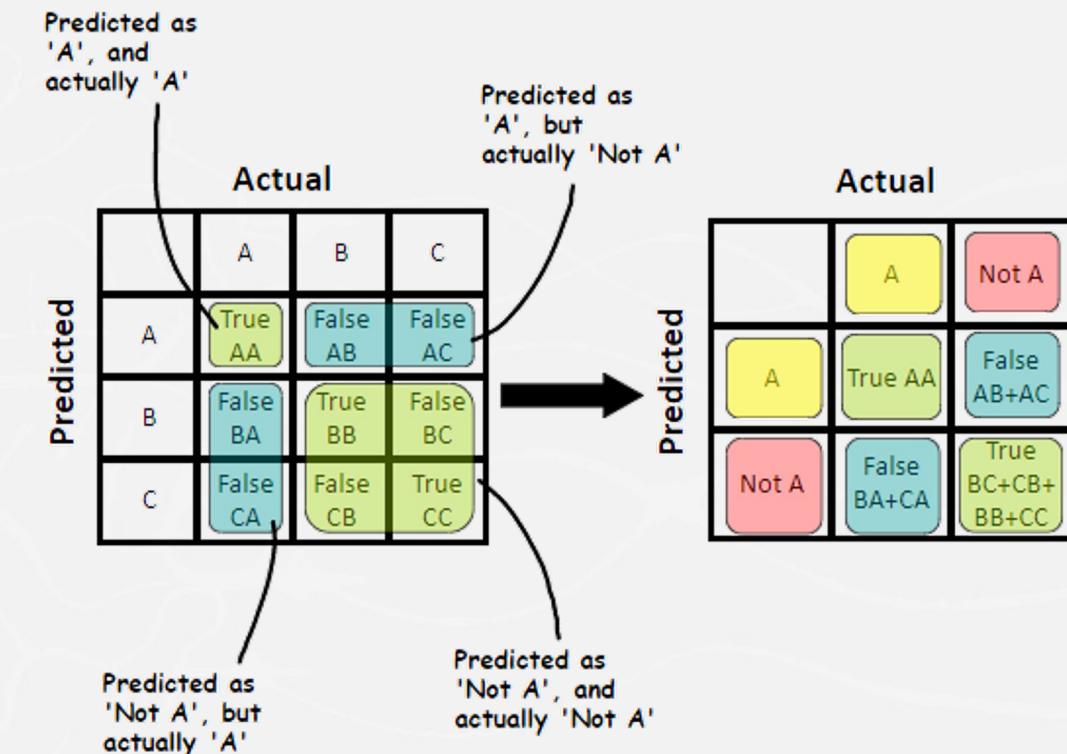
## 2- Examples of machine learning applications

		Predicted Class		
		Positive	Negative	
Actual Class	Positive	True Positive (TP)	False Negative (FN) <b>Type II Error</b>	Sensitivity $\frac{TP}{(TP + FN)}$
	Negative	False Positive (FP) <b>Type I Error</b>	True Negative (TN)	Specificity $\frac{TN}{(TN + FP)}$
		Precision $\frac{TP}{(TP + FP)}$	Negative Predictive Value $\frac{TN}{(TN + FN)}$	Accuracy $\frac{TP + TN}{(TP + TN + FP + FN)}$

id	color
1	red
2	blue
3	green
4	blue

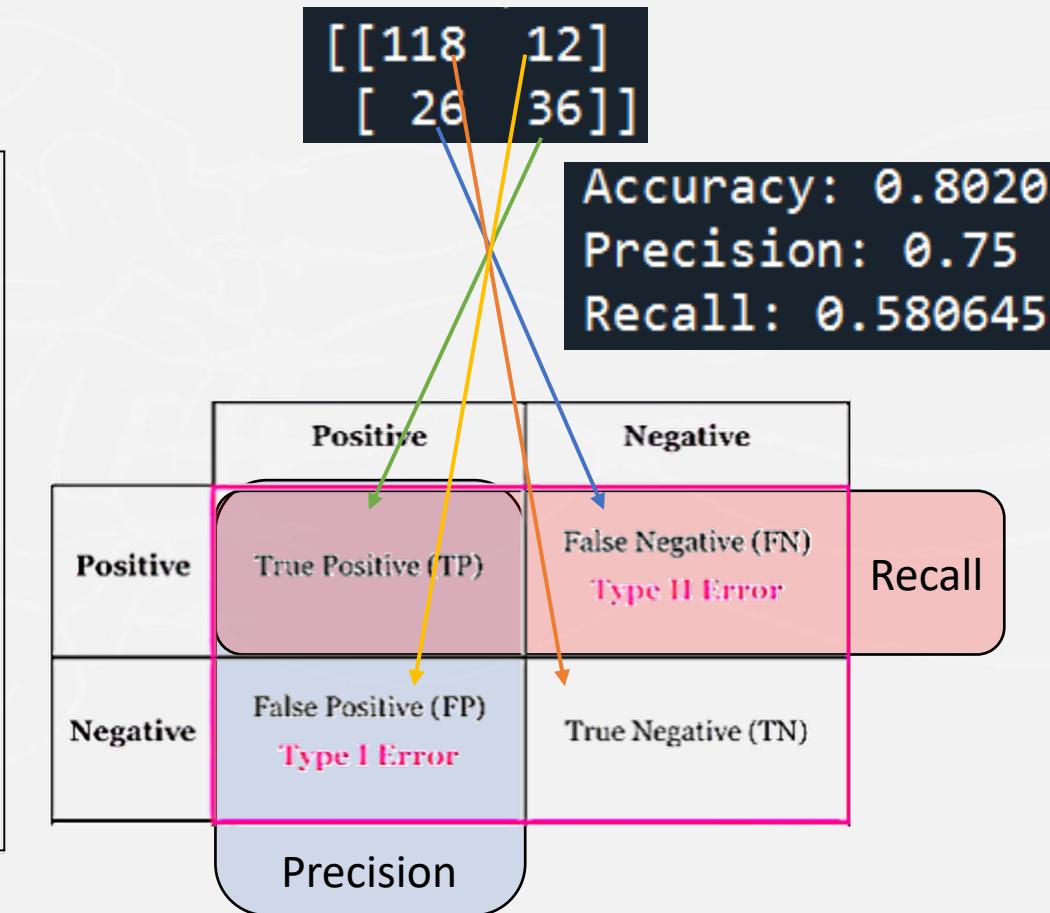
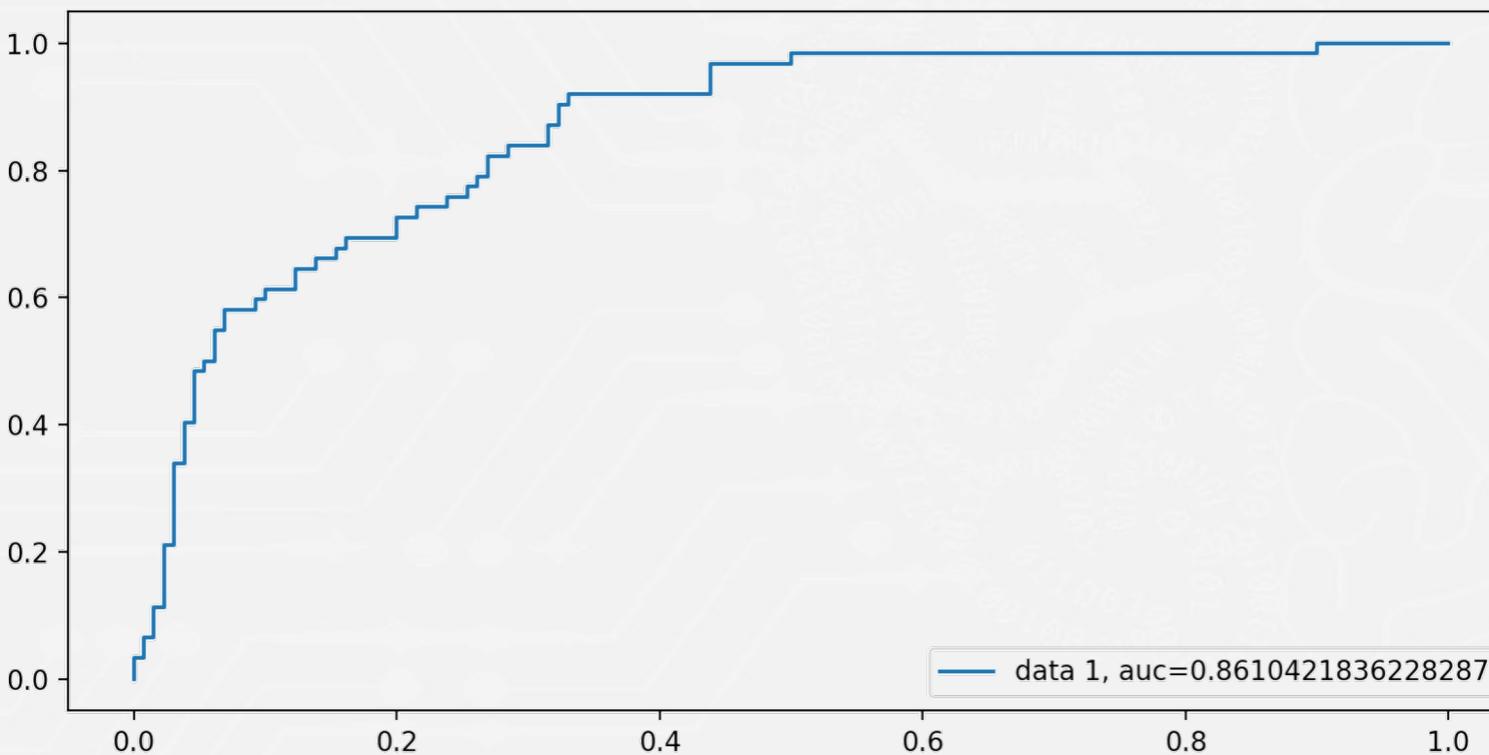
One Hot Encoding →

id	color_red	color_blue	color_green
1	1	0	0
2	0	1	0
3	0	0	1
4	0	1	0



# 2- Examples of machine learning applications

## Classification - Logistic regression

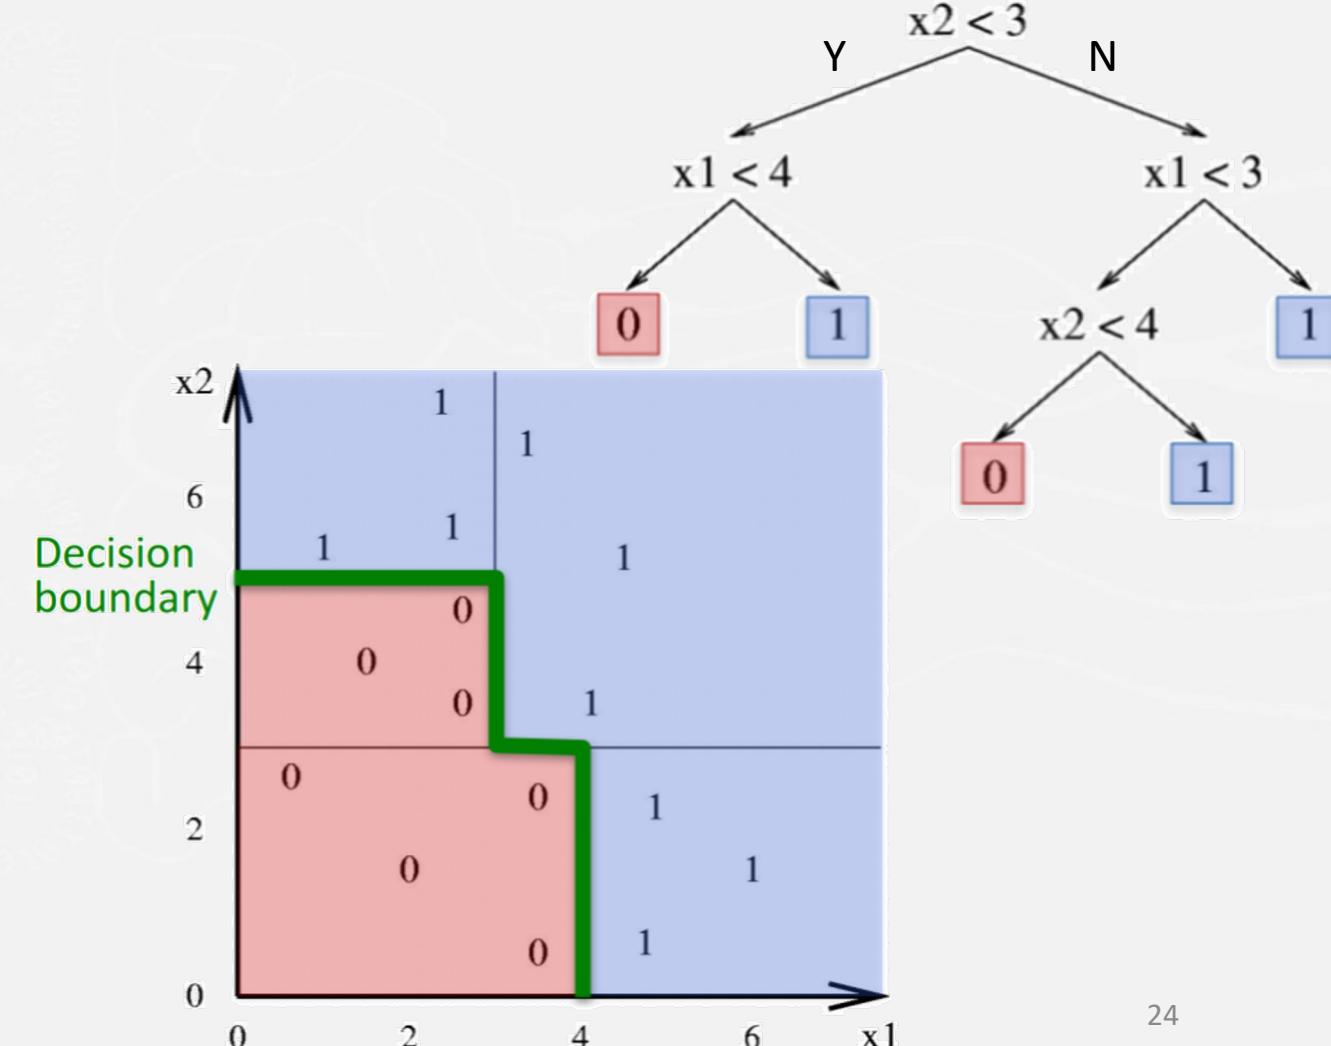


## 2- Examples of machine learning applications

### Classification - Decision tree

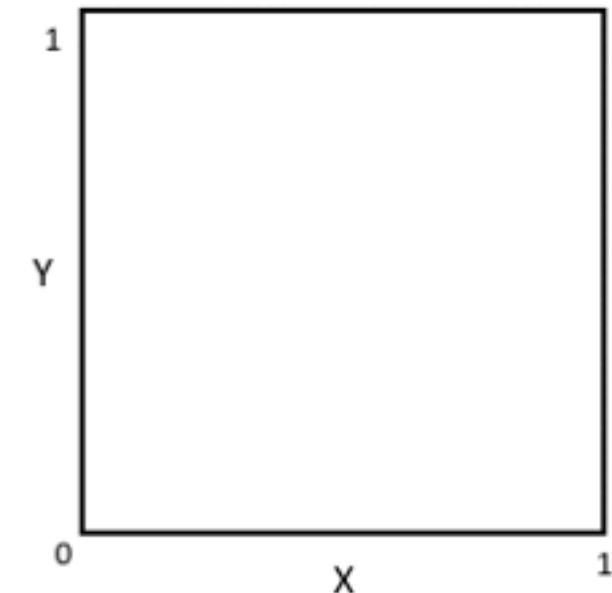
Divide the feature space into axis-parallel (**hyper**)rectangles

Each rectangular region is labeled with one label or a probability distribution over labels



# 2- Examples of machine learning applications

## Classification - Decision tree



## 2- Examples of machine learning applications

### Iris flower species prediction

### Fisher's Iris Dataset

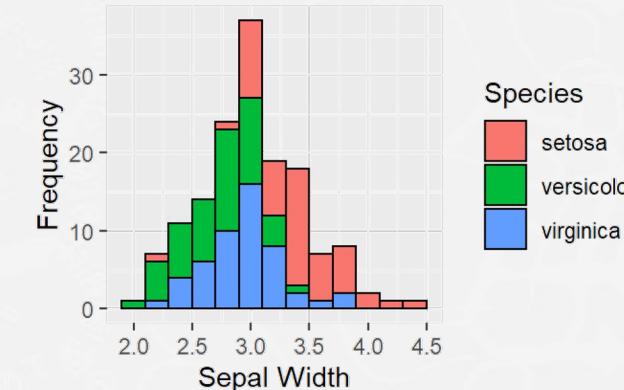


Iris Versicolor

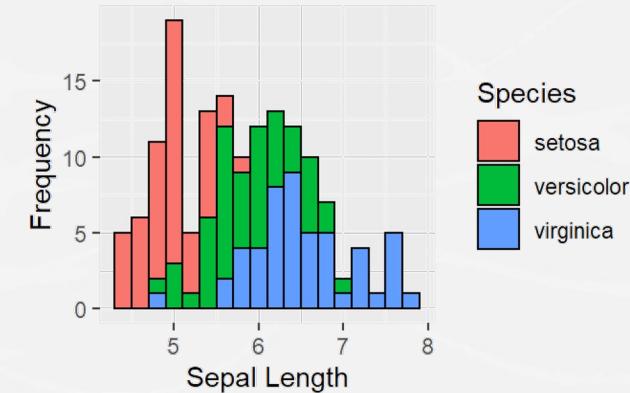
Iris Setosa

Iris Virginica

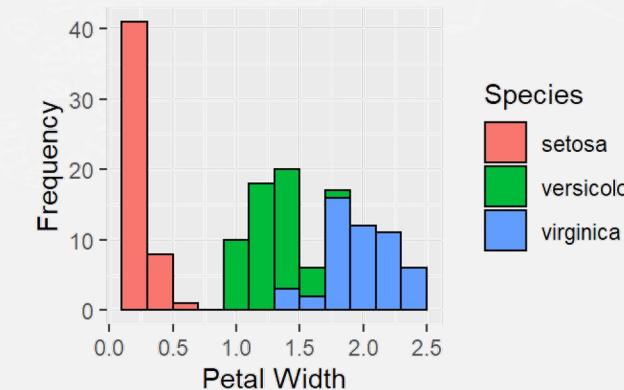
Histogram of Sepal Width



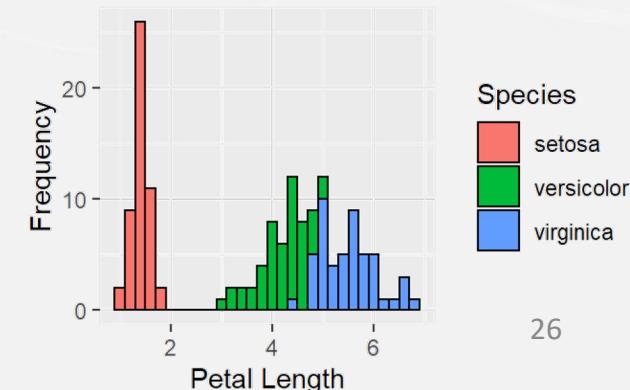
Histogram of Sepal Length



Histogram of Petal Width



Histogram of Petal Length



# Classification - Decision tree

```

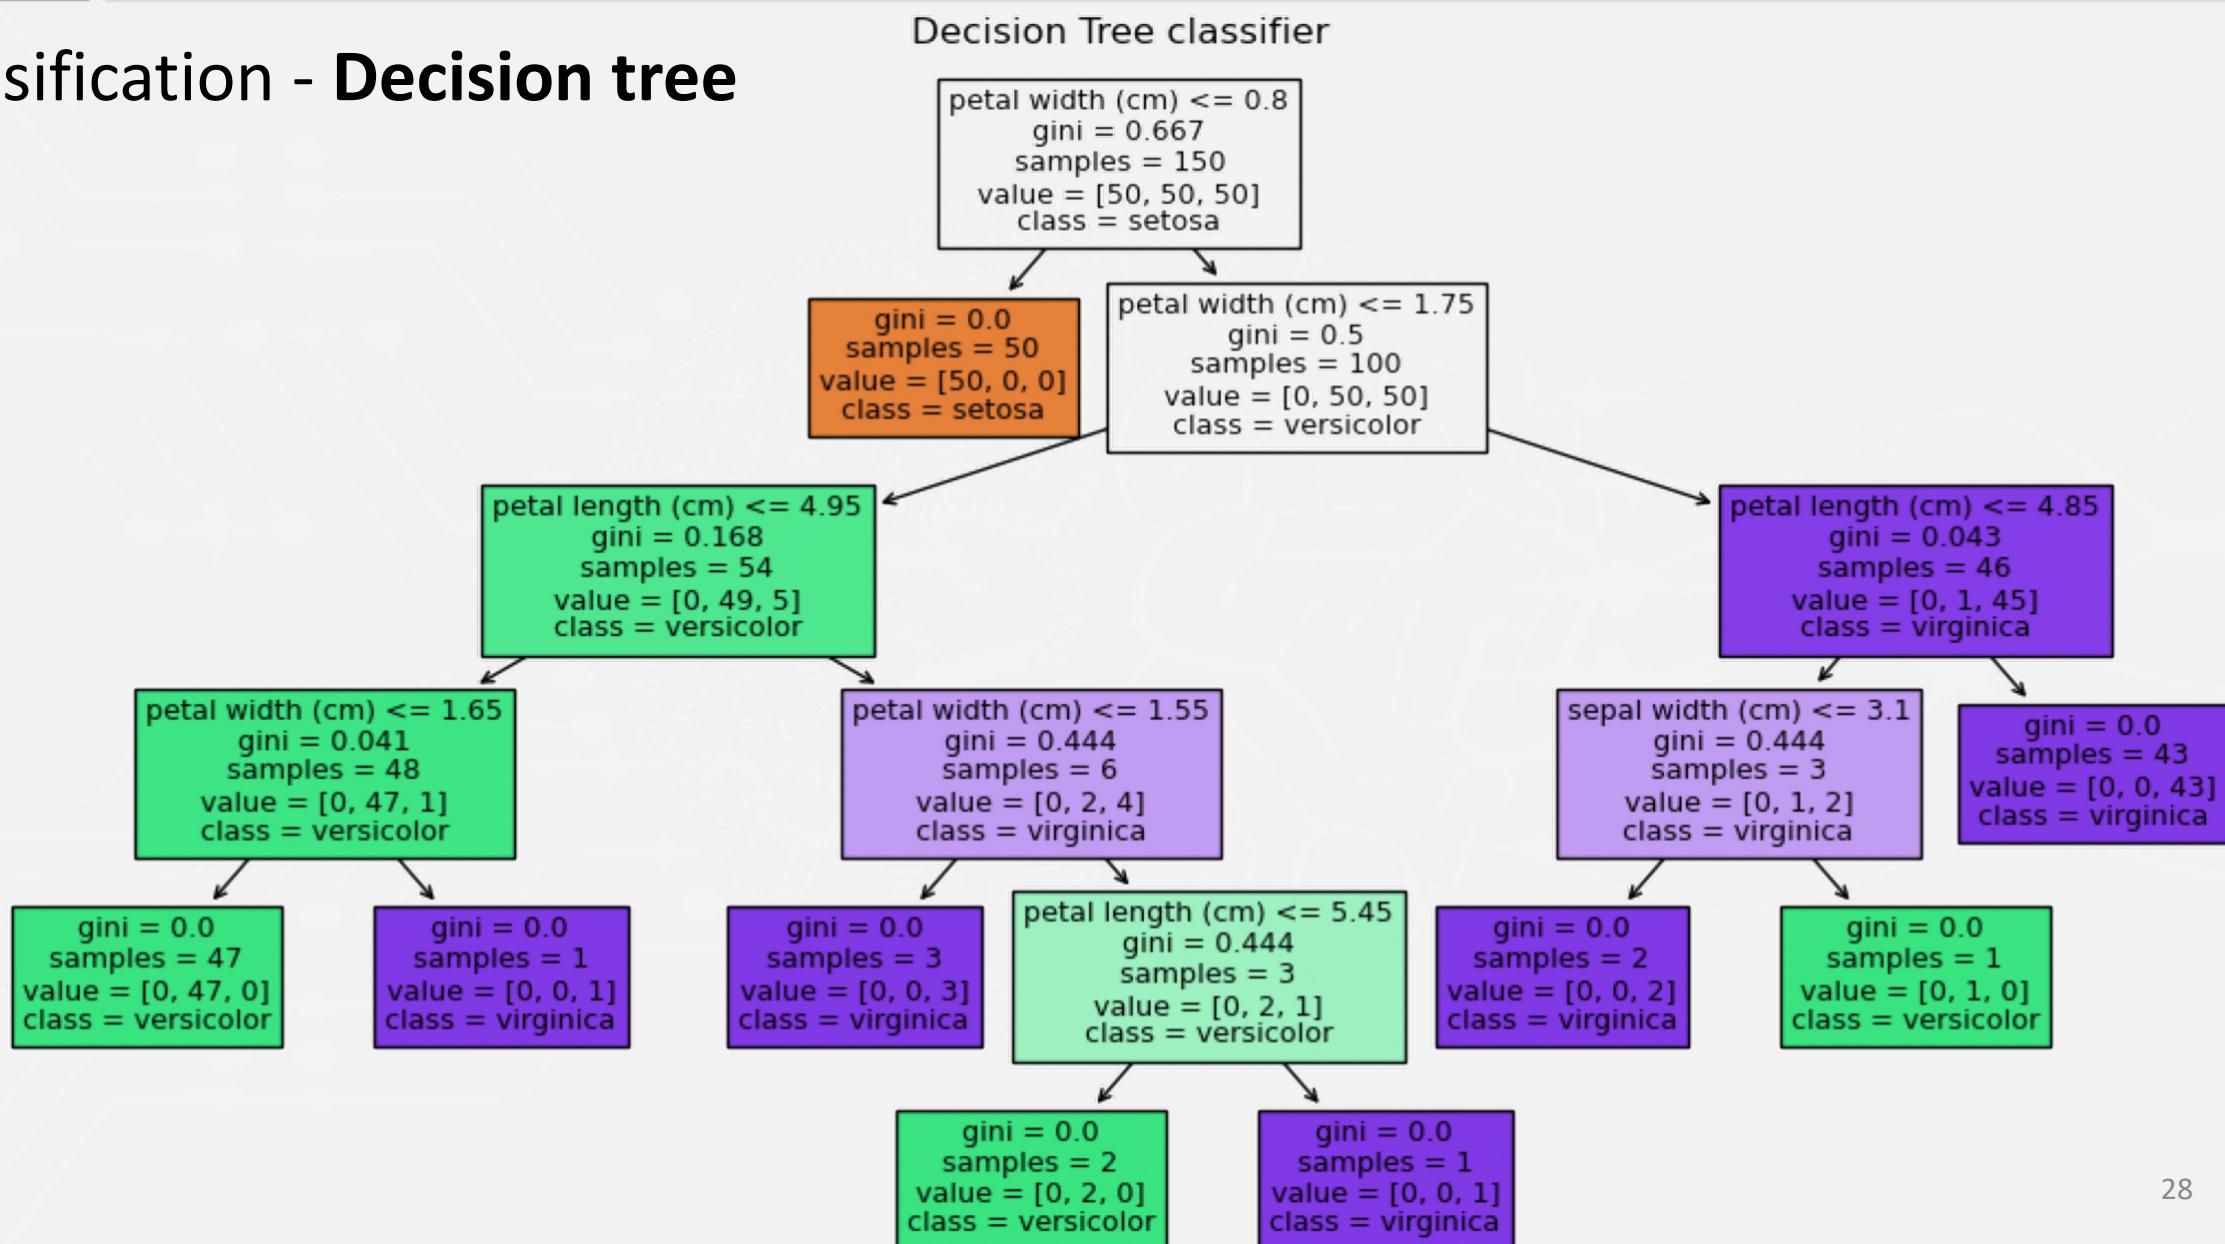
1     """
2     Import libraries
3     """
4     import numpy as np
5     import matplotlib.pyplot as plt
6     from sklearn.datasets import load_iris
7     from sklearn.tree import DecisionTreeClassifier
8     from sklearn.model_selection import train_test_split, cross_val_score
9     from sklearn.tree import plot_tree
10    """
11    Import the dataset
12    """
13    iris = load_iris()
14    X = iris.data
15    y = iris.target
16    """
17    Define the model
18    """
19    clf = DecisionTreeClassifier()
20    """
21    Train and validate the model using cross-validation
22    """
23    cv_scores = cross_val_score(clf, X, y, cv=5)
24    print("Cross-validation scores:", cv_scores)
25    print("Mean CV accuracy:", np.mean(cv_scores))
26    """
27    Train the model on the entire training set
28    """
29    clf.fit(X, y)
30    """
31    Plot the decision tree
32    """
33    plt.figure(figsize=(15, 10))
34    plot_tree(clf, filled=True, feature_names=iris.feature_names,
35              class_names=list(iris.target_names))
36    plt.title("Decision Tree classifier")
37    plt.show()

```



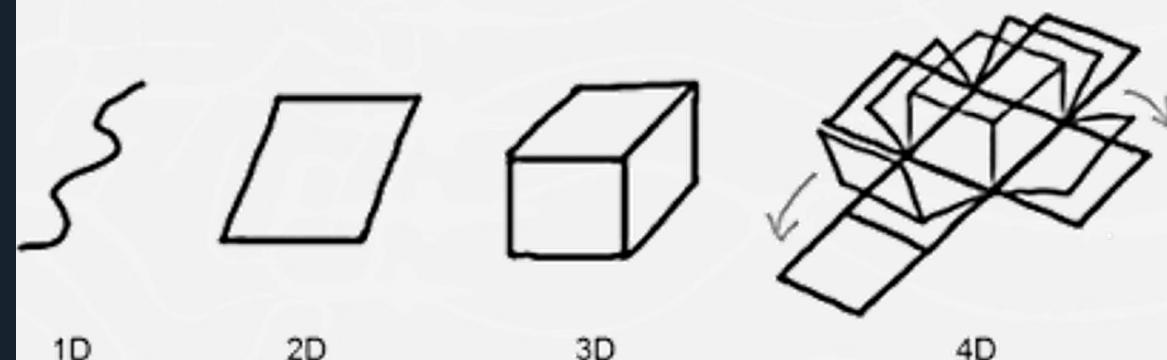
Cross-validation scores: [0.96666667 0.96666667 0.96666667 0.96666667 1.]  
 Mean CV accuracy: 0.9600000000000002

# Classification - Decision tree

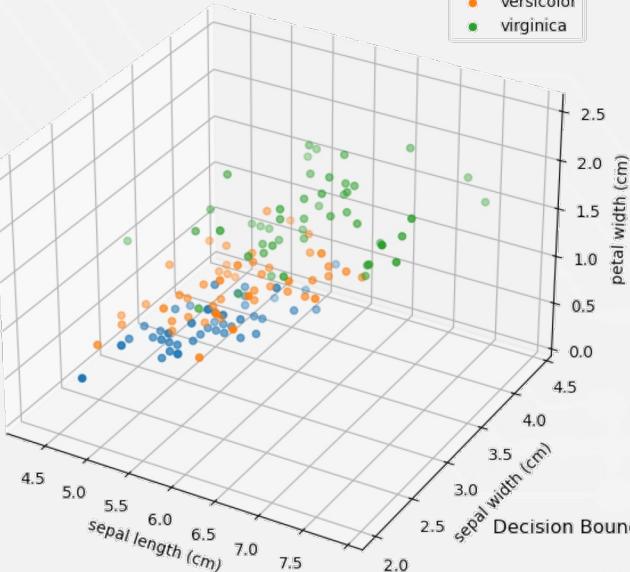


# Classification - Decision tree

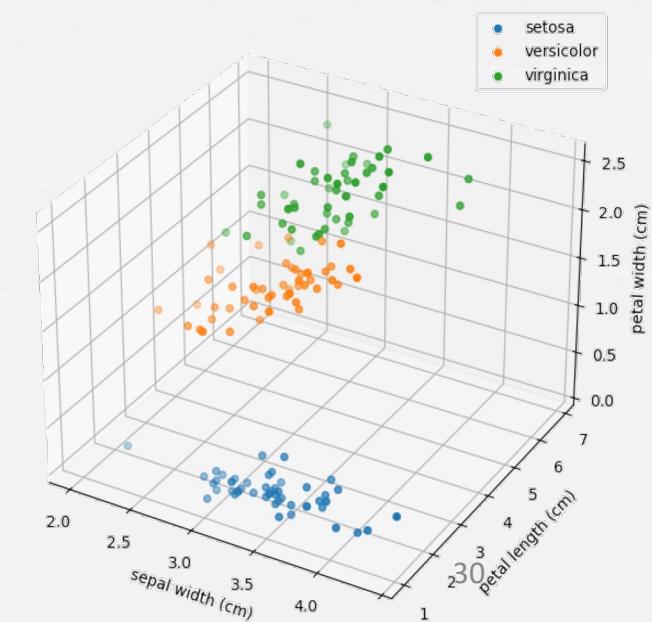
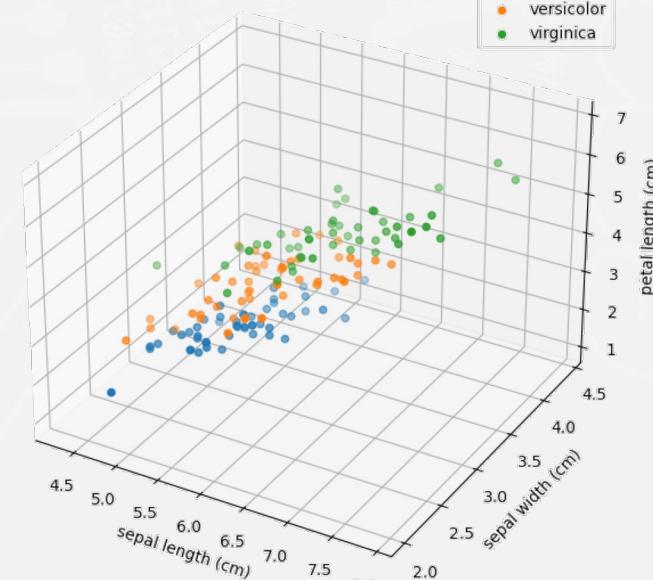
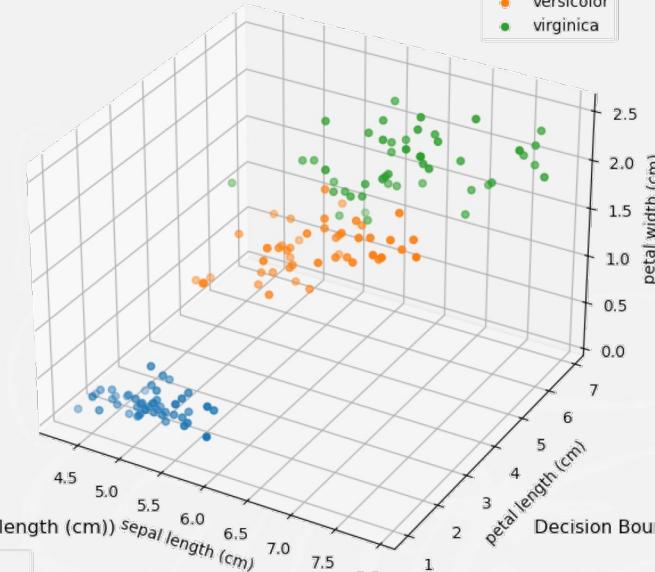
```
38 """
39 Show the decision boundaries by first defining the feature triplets for plotting
40 """
41 feature_triplets = [(0, 1, 2), (0, 1, 3), (0, 2, 3), (1, 2, 3)]
42 triplet_names = [
43     (iris.feature_names[i], iris.feature_names[j], iris.feature_names[k])
44     for i, j, k in feature_triplets
45 ]
46 """
47 Create separate 3D plots for each feature triplet
48 """
49 for i, (x_feature, y_feature, z_feature) in enumerate(triplet_names):
50     fig = plt.figure(figsize=(10, 8))
51     ax = fig.add_subplot(111, projection='3d')
52     x_index, y_index, z_index = feature_triplets[i]
53     X_triplet = X[:, [x_index, y_index, z_index]]
54     #plot the data points
55     for j, target_name in enumerate(iris.target_names):
56         ax.scatter(
57             X_triplet[y == j, 0],
58             X_triplet[y == j, 1],
59             X_triplet[y == j, 2],
60             label=target_name
61         )
62     #set the axis labels
63     ax.set_xlabel(x_feature)
64     ax.set_ylabel(y_feature)
65     ax.set_zlabel(z_feature)
66     ax.set_title(f"Decision Boundaries in 3D ({x_feature}, {y_feature}, {z_feature})")
67     plt.legend()
68     #show the plot
69     plt.show()
```



Decision Boundaries in 3D (sepal length (cm), sepal width (cm), petal width (cm))



Decision Boundaries in 3D (sepal length (cm), petal length (cm), petal width (cm))



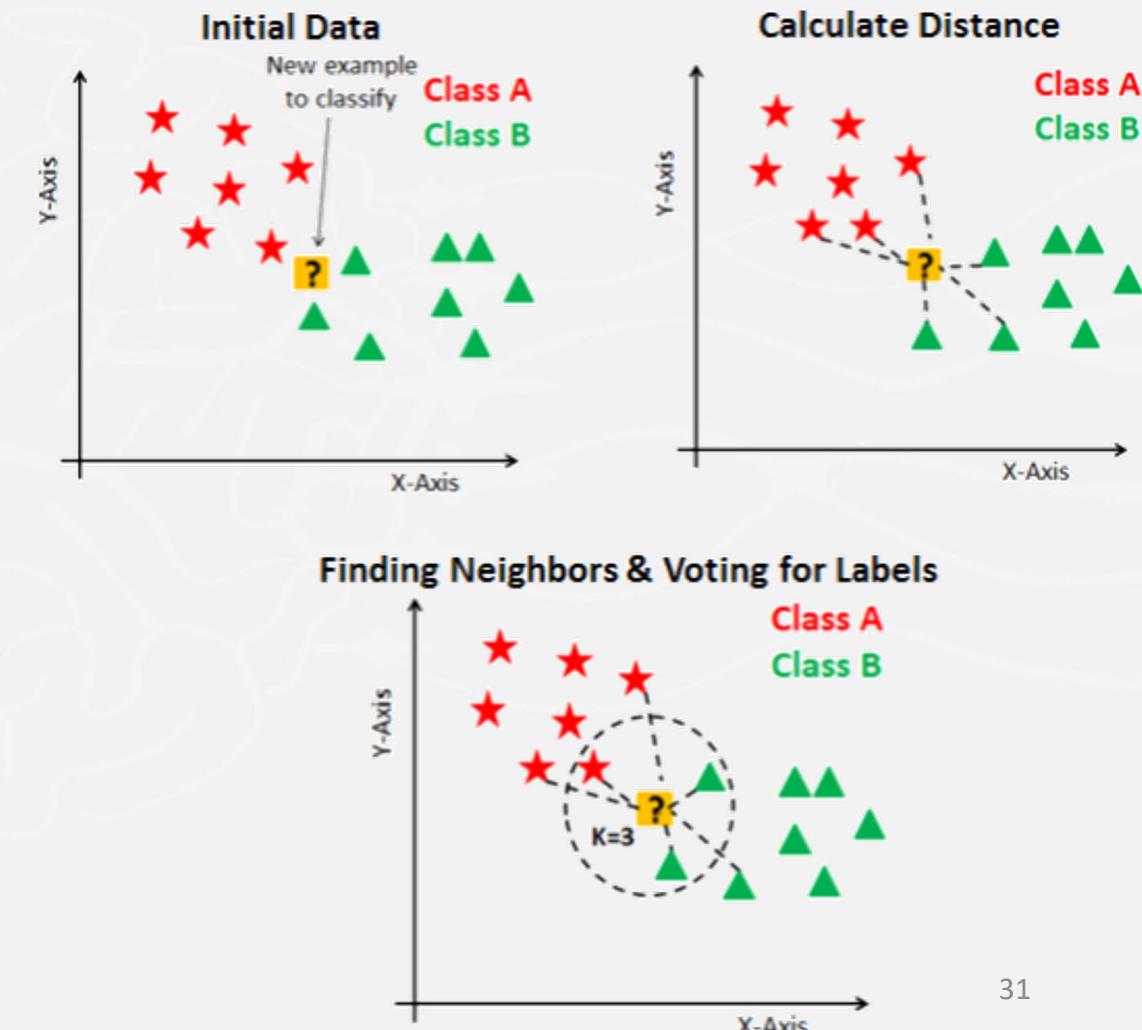
## 2- Examples of machine learning applications

### Classification - K Nearest Neighbor (kNN)

Calculates the distance of a new data point to all other training data points.

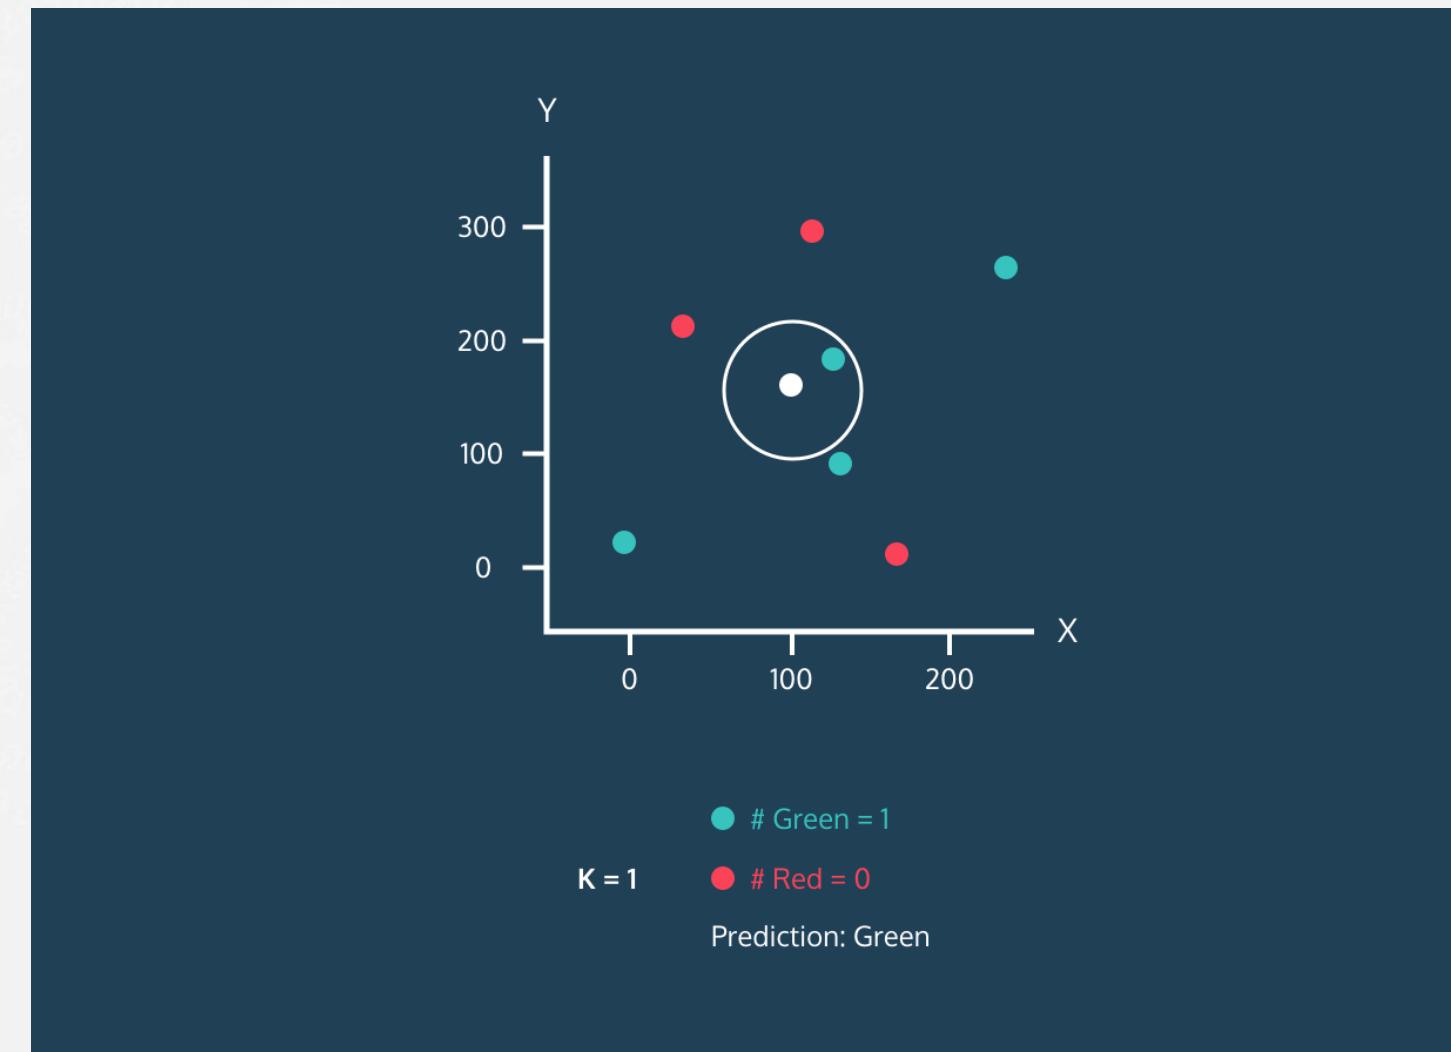
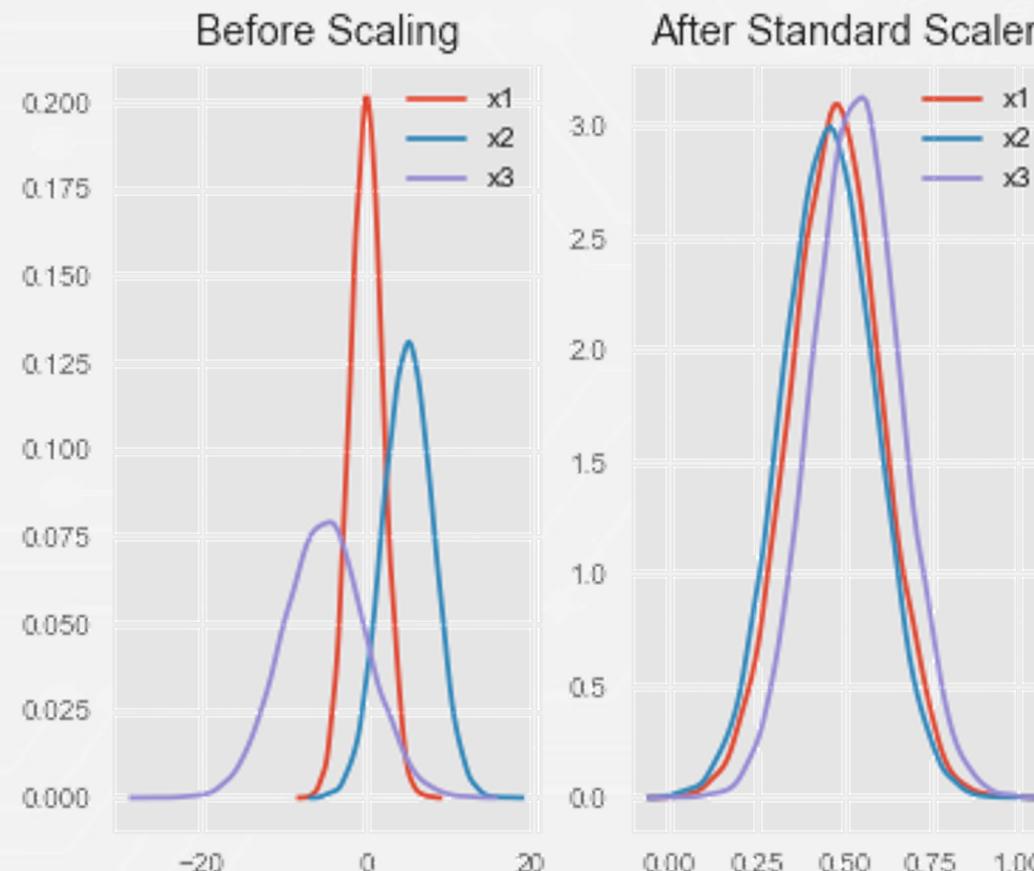
Selects the K-nearest data points.

Assigns the data point to the class to which the majority of the K data points belong.



# 2- Examples of machine learning applications

## Classification - kNN



# 2- Examples of machine learning applications

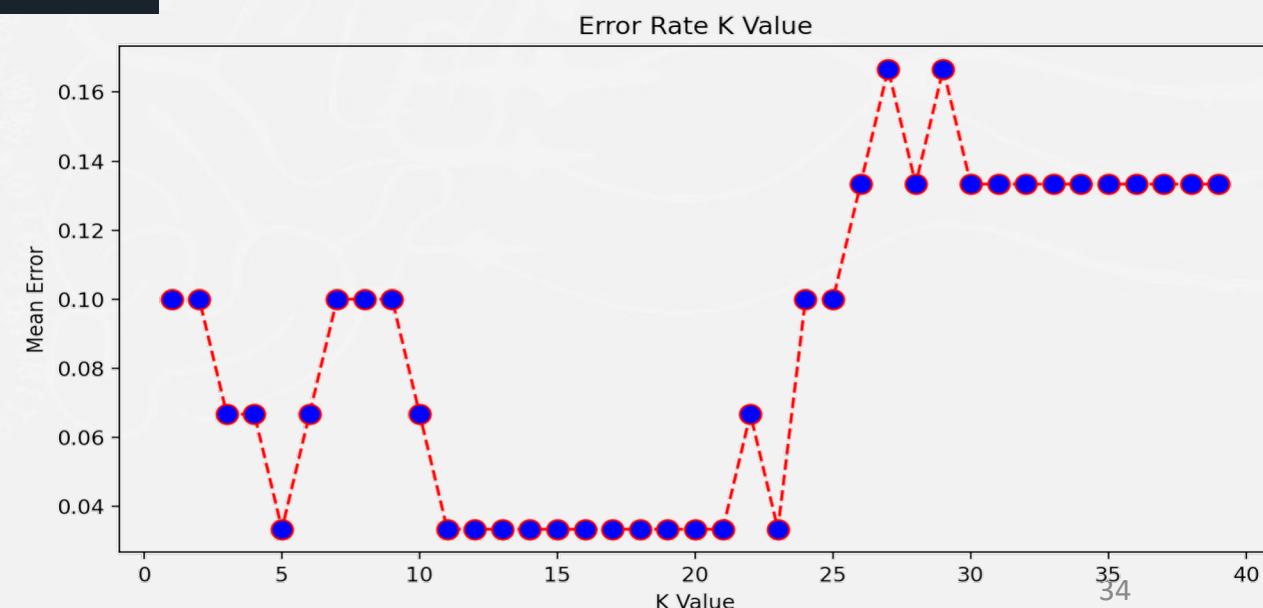
## Classification - kNN

```
1      """
2  Importing Libraries
3      """
4  import numpy as np
5  import matplotlib.pyplot as plt
6  import pandas as pd
7      """
8  Importing the dataset
9      """
10 url="https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
11 # Assign colum names to the dataset
12 names=['sepal-length', 'sepal-width', 'petal-length', 'petal-width', 'Class']
13 # Read dataset to pandas datafram
14 dataset = pd.read_csv(url, names=names)
15 # Check the data
16 dataset.head()
17 """
18 Preprocessing
19 """
20 X = dataset.iloc[:, :-1].values
21 y = dataset.iloc[:, 4].values
22 """
23 Create training and test splits,
24 """
25 from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20)
```

```
27      """
28 Feature scaling
29 """
30 from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
"""
36 Training and predictions
37 """
38 from sklearn.neighbors import KNeighborsClassifier
classifier = KNeighborsClassifier(n_neighbors=5)
classifier.fit(X_train, y_train)
y_pred = classifier.predict(X_test)
"""
43 Evaluating the algorithm
44 """
45 from sklearn.metrics import classification_report, confusion_matrix
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
48
49     """Comparing error rate with the K value"""
50
51     error = []
52     # Calculating error for K values between 1 and 40
53     for i in range(1, 40):
54         knn = KNeighborsClassifier(n_neighbors=i)
55         knn.fit(X_train, y_train)
56         pred_i = knn.predict(X_test)
57         error.append(np.mean(pred_i != y_test))
58     plt.figure(figsize=(12, 6))
59     plt.plot(range(1, 40), error, color='red', linestyle='dashed', marker='o',
60             markerfacecolor='blue', markersize=10)
61     plt.title('Error Rate K Value')
62     plt.xlabel('K Value')
63     plt.ylabel('Mean Error')
```

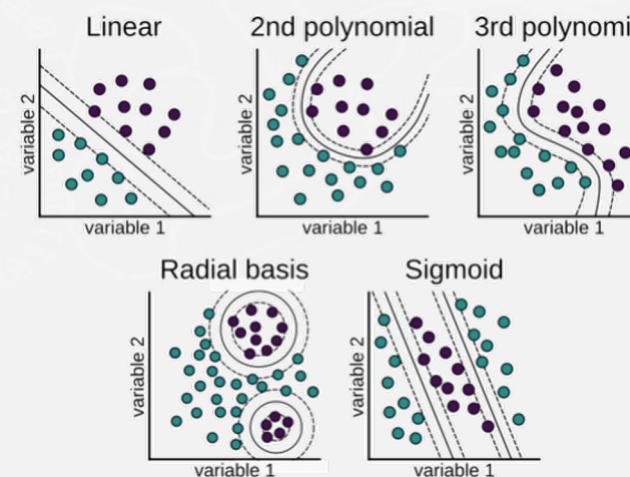
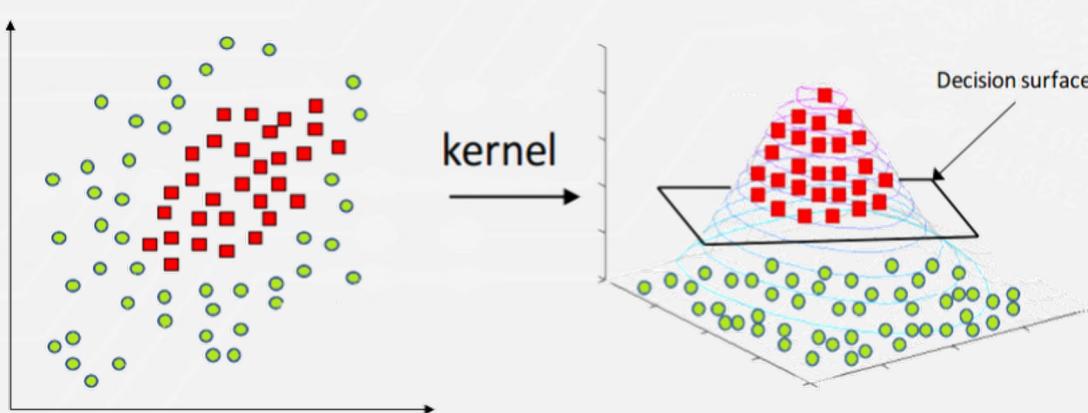
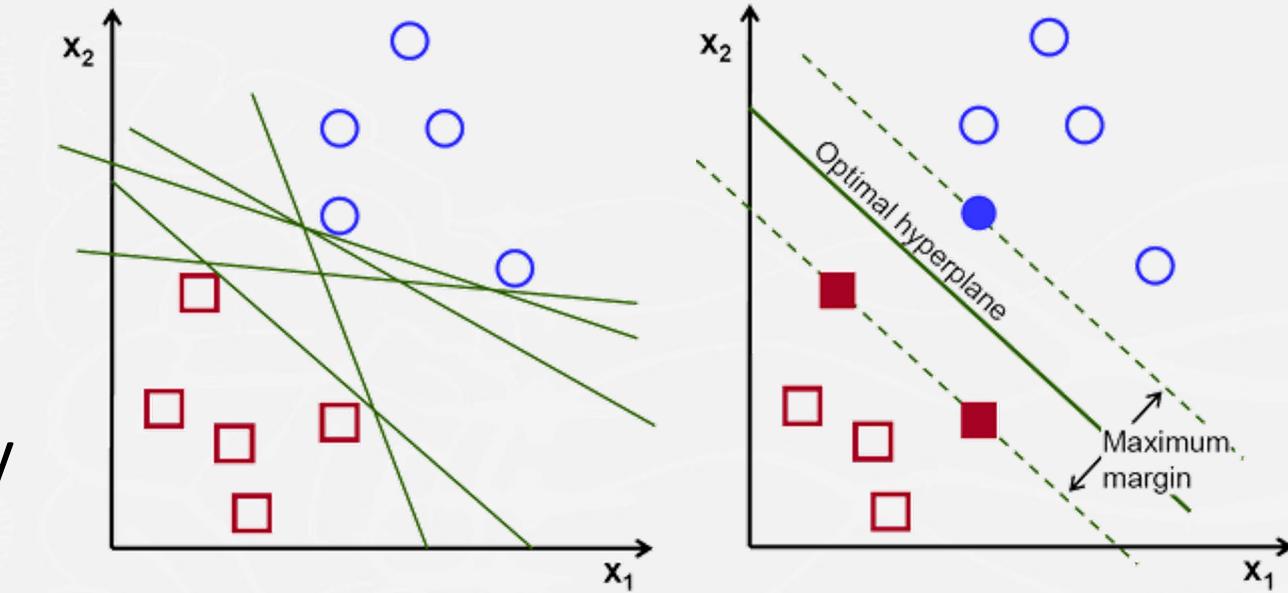
[[ 8 0 0]
[ 0 10 0]
[ 0 1 11]]
precision recall f1-score support
Iris-setosa 1.00 1.00 1.00 8
Iris-versicolor 0.91 1.00 0.95 10
Iris-virginica 1.00 0.92 0.96 12
accuracy 0.97 0.97 0.97 30
macro avg 0.97 0.97 0.97 30
weighted avg 0.97 0.97 0.97 30



## 2- Examples of machine learning applications

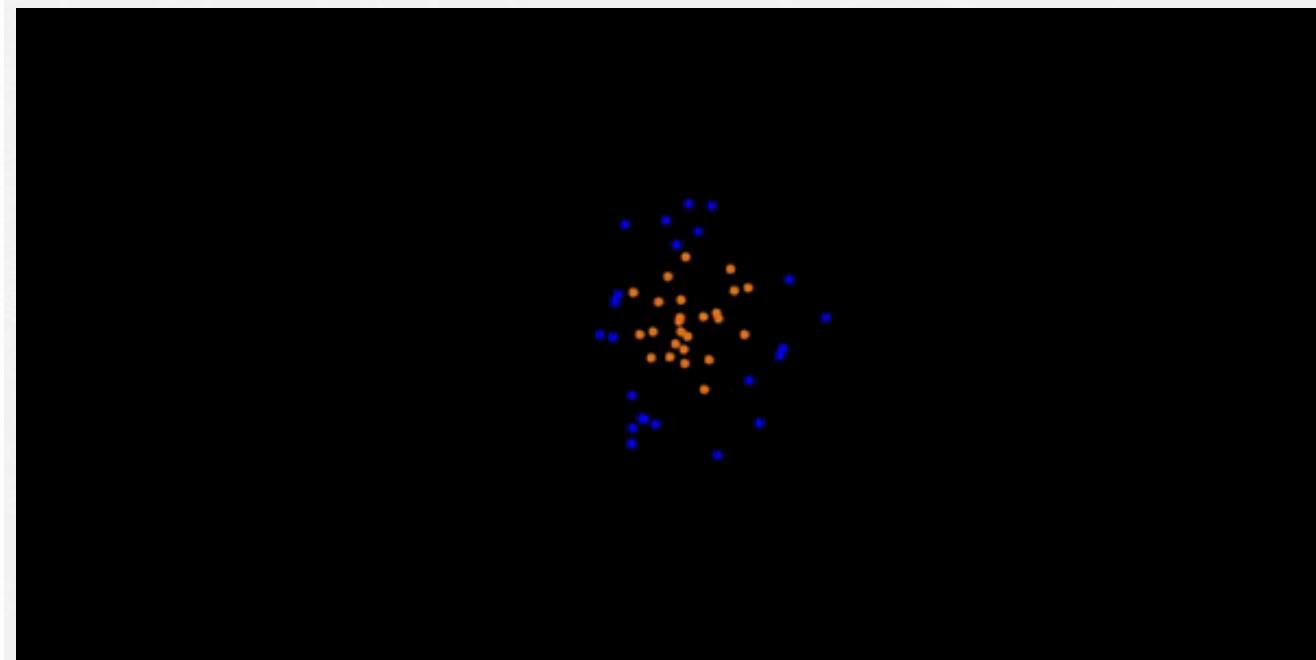
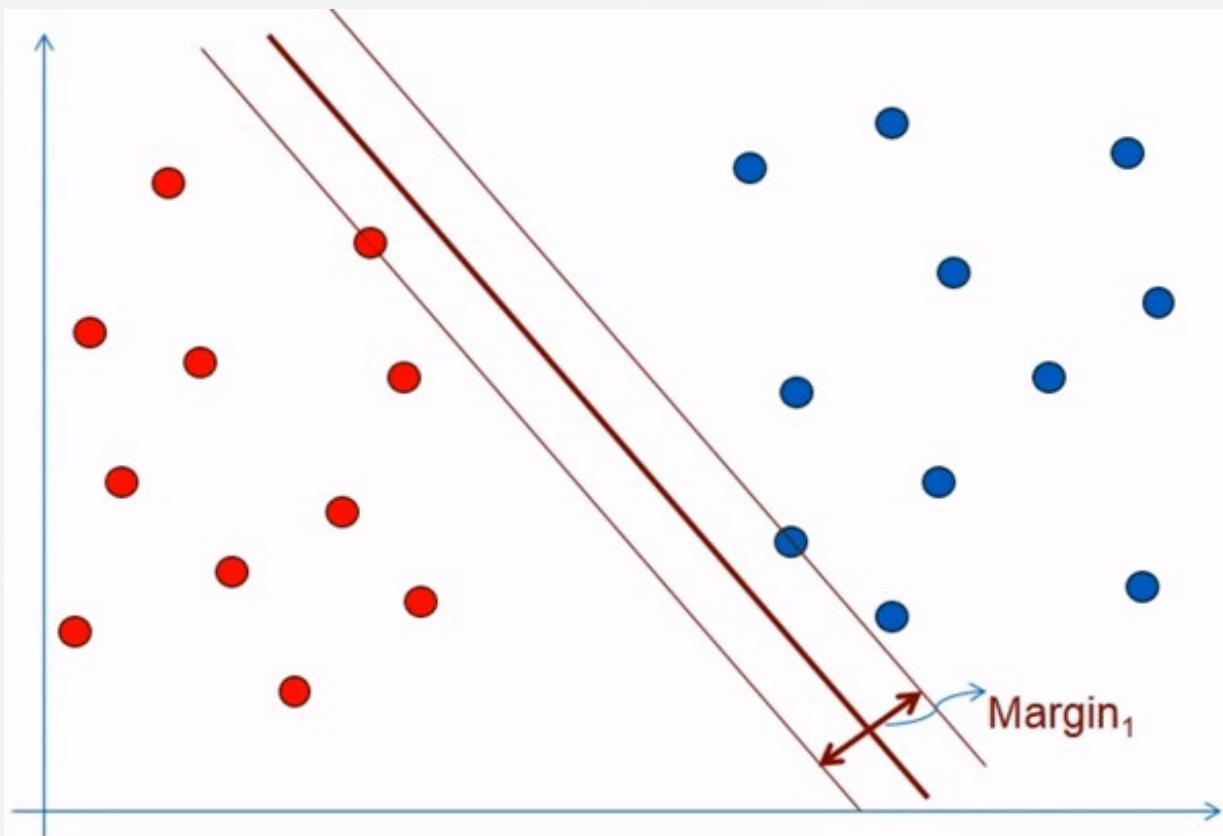
### Classification - Support Vector Machine (SVM)

Find a hyperplane in an N-dimensional space (N — the number of features) that distinctly classifies the data points.



## 2- Examples of machine learning applications

### Classification - SVM



# Classification - SVM

```
1      """
2      Importing Libraries
3      """
4      import pandas as pd
5      """
6      Importing the dataset
7      """
8      url="https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
9      # Assign colum names to the dataset
10     colnames=['sepal-length', 'sepal-width', 'petal-length', 'petal-width', 'Class']
11     # Read dataset to pandas dataframe
12     irisdata = pd.read_csv(url, names=colnames)
13     """
14     Preprocessing
15     """
16     X = irisdata.drop('Class', axis=1)
17     y = irisdata['Class']
18     """
19     Train Test Split
20     """
21     from sklearn.model_selection import train_test_split
22     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30)
```

```
23    """
24    Training the algorithm with Gaussian Kernel
25    """
26    from sklearn.svm import SVC
27    svclassifier = SVC(kernel='rbf')
28    svclassifier.fit(X_train, y_train)
29    """
30    Prediction and evaluation
31    """
32    y_pred = svclassifier.predict(X_test)
33    from sklearn.metrics import classification_report, confusion_matrix
34    print(confusion_matrix(y_test, y_pred))
35    print(classification_report(y_test, y_pred))
```

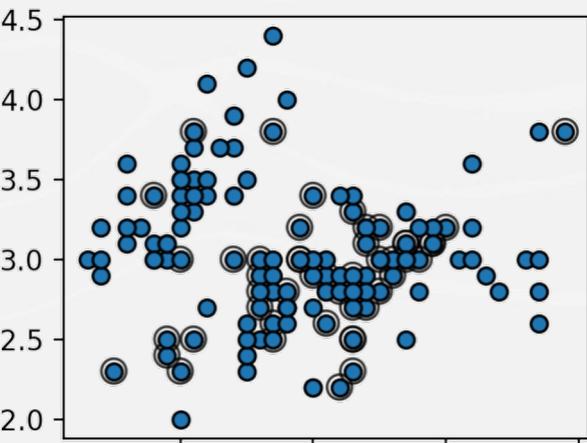
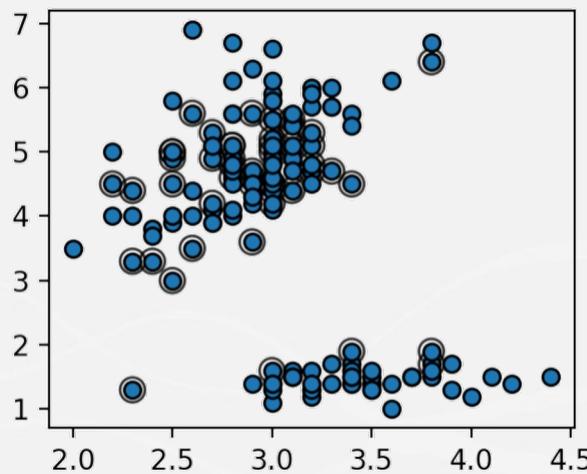
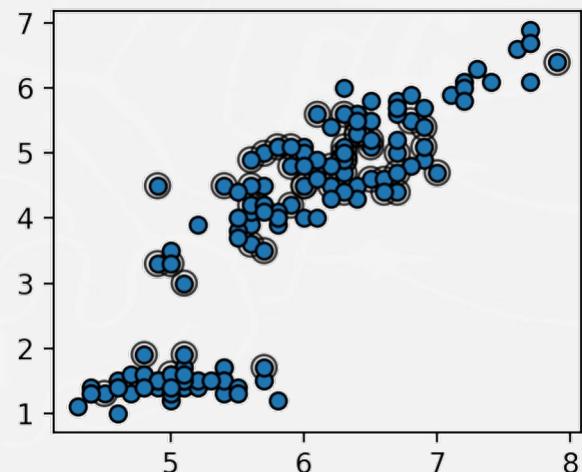
[[18 0 0]				
[ 0 6 0]				
[ 0 0 21]]				
	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	18
Iris-versicolor	1.00	1.00	1.00	6
Iris-virginica	1.00	1.00	1.00	21
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45

## Examples of support vectors for three of the four features

```

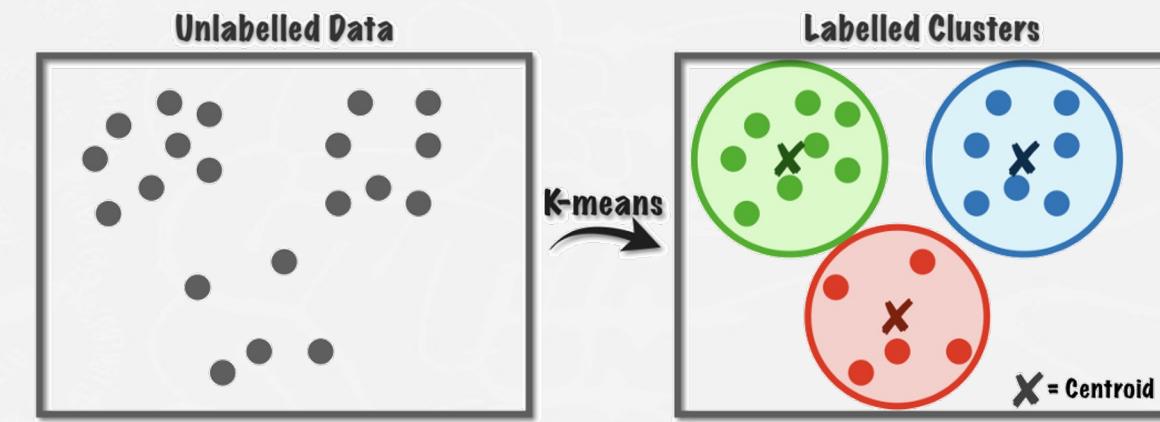
36
37     """
38     Plot the suport vectors
39     """
40
41     import matplotlib.pyplot as plt
42     plt.figure(0)
43     plt.scatter(svclassifier.support_vectors_[:, 0],
44                 svclassifier.support_vectors_[:, 1], s=80,
45                 facecolors='none', edgecolors='k', cmap=plt.cm.coolwarm, alpha=0.8)
46     plt.scatter(X.to_numpy()[:, 0], X.to_numpy()[:, 1], cmap=plt.cm.coolwarm,
47                 edgecolors='k')
48     plt.figure(1)
49     plt.scatter(svclassifier.support_vectors_[:, 0],
50                 svclassifier.support_vectors_[:, 2], s=80,
51                 facecolors='none', edgecolors='k', cmap=plt.cm.coolwarm, alpha=0.8)
52     plt.scatter(X.to_numpy()[:, 0], X.to_numpy()[:, 2], cmap=plt.cm.coolwarm,
53                 edgecolors='k')
54     plt.figure(2)
55     plt.scatter(svclassifier.support_vectors_[:, 1],
56                 svclassifier.support_vectors_[:, 2], s=80,
57                 facecolors='none', edgecolors='k', cmap=plt.cm.coolwarm, alpha=0.8)
58     plt.scatter(X.to_numpy()[:, 1], X.to_numpy()[:, 2], cmap=plt.cm.coolwarm,
59                 edgecolors='k')

```

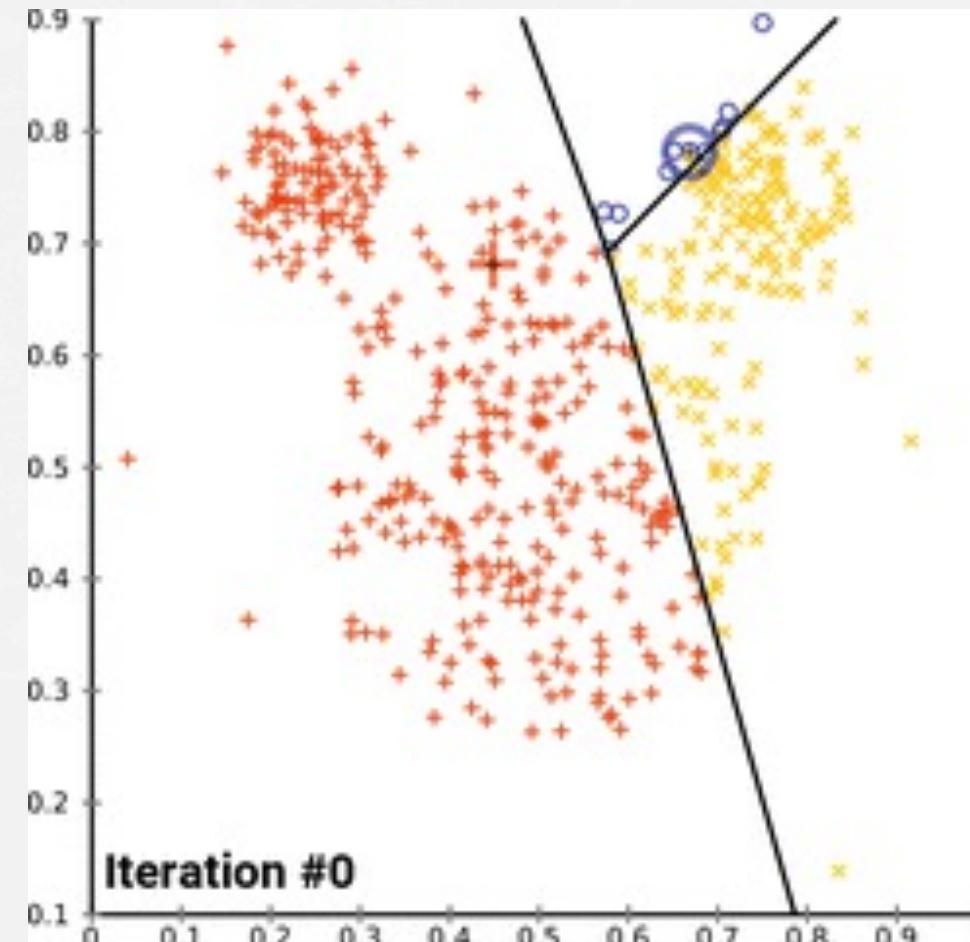


## Clustering - K-means

Identifies K number of **centroids**, and then allocates every data point to the nearest cluster, while keeping the centroids as small as possible.



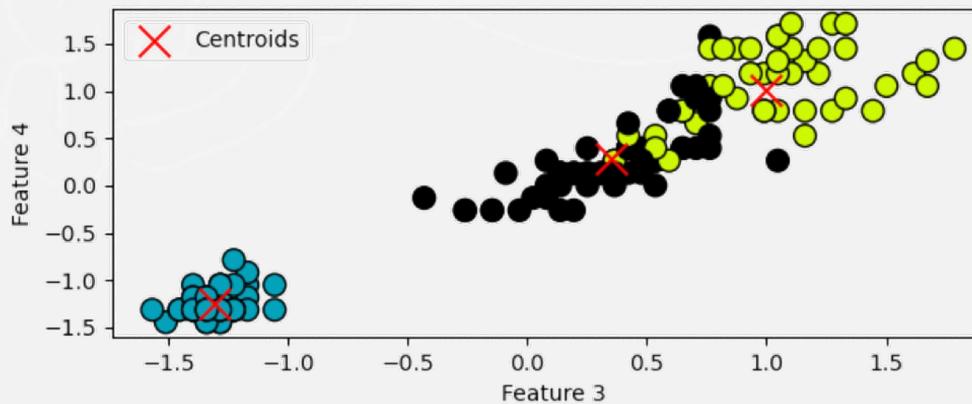
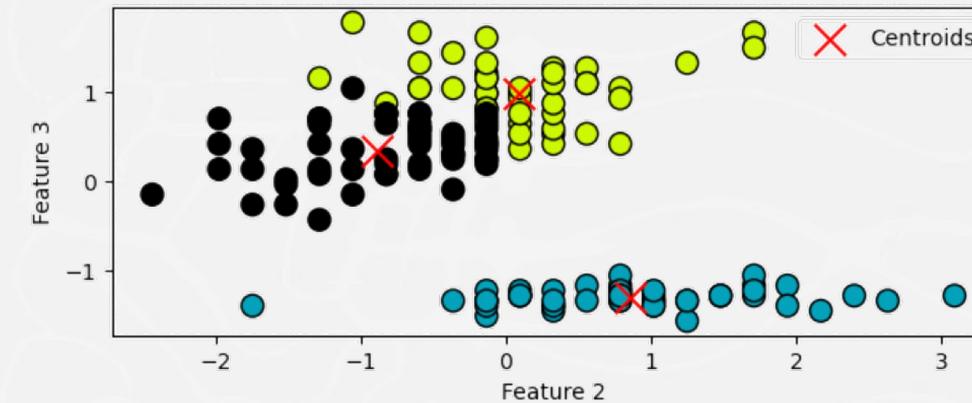
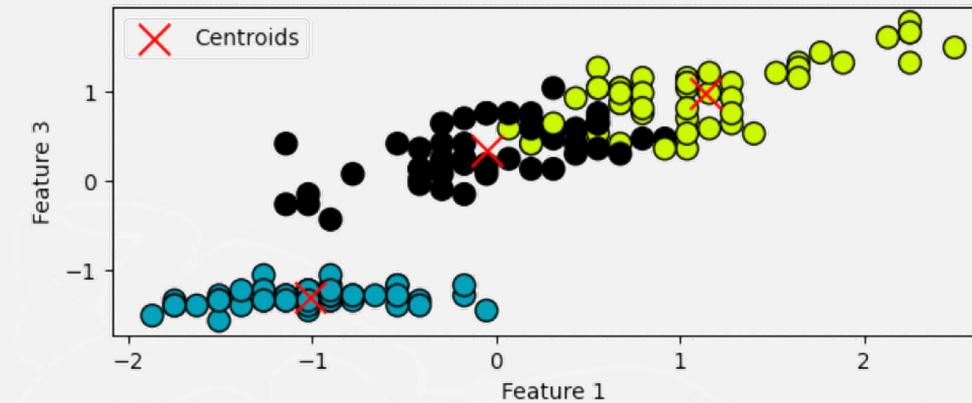
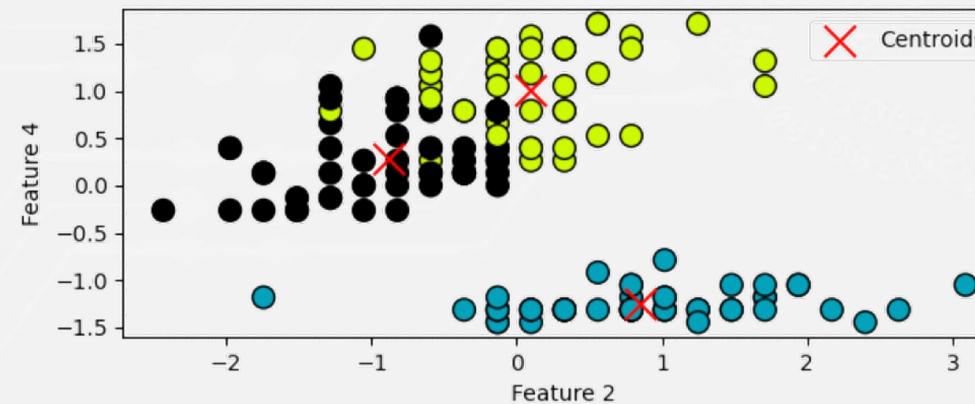
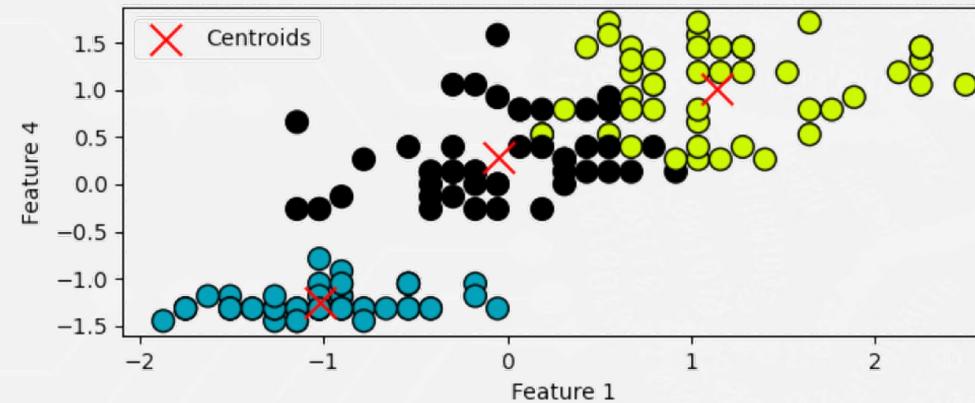
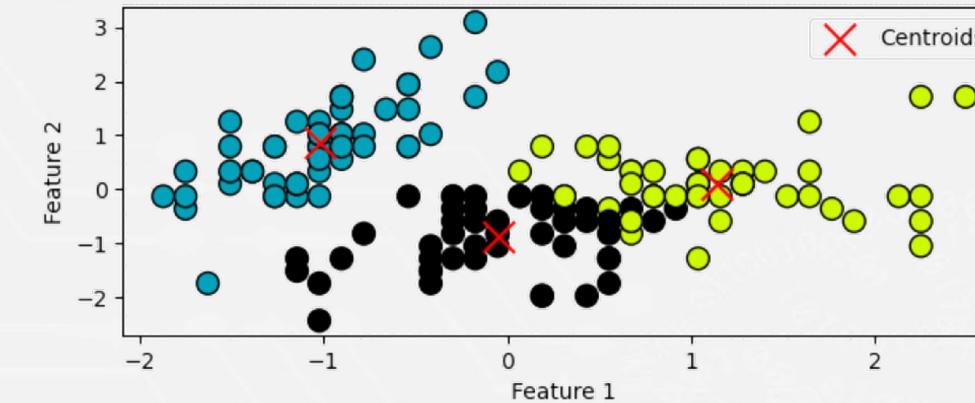
## Clustering - K-means



# Clustering - K-means

```
1  """
2  Import libraries
3  """
4
5  import numpy as np
6  import matplotlib.pyplot as plt
7  from sklearn.datasets import load_iris
8  from sklearn.cluster import KMeans
9  from sklearn.metrics import silhouette_samples, silhouette_score
10 from sklearn.preprocessing import StandardScaler
11 from mpl_toolkits.mplot3d import Axes3D
12 import matplotlib.cm as cm
13 from itertools import combinations
14 """
15 Import the dataset
16 """
17 iris = load_iris()
18 X = iris.data
19 """
20 Standardize the features
21 """
22 scaler = StandardScaler()
23 X_scaled = scaler.fit_transform(X)
24 """
25 Set the number of clusters
26 """
27 n_clusters = 3
28 """
29 Perform k-means clustering
30 """
31 kmeans = KMeans(n_clusters=n_clusters, random_state=42)
32 cluster_labels = kmeans.fit_predict(X_scaled)
```

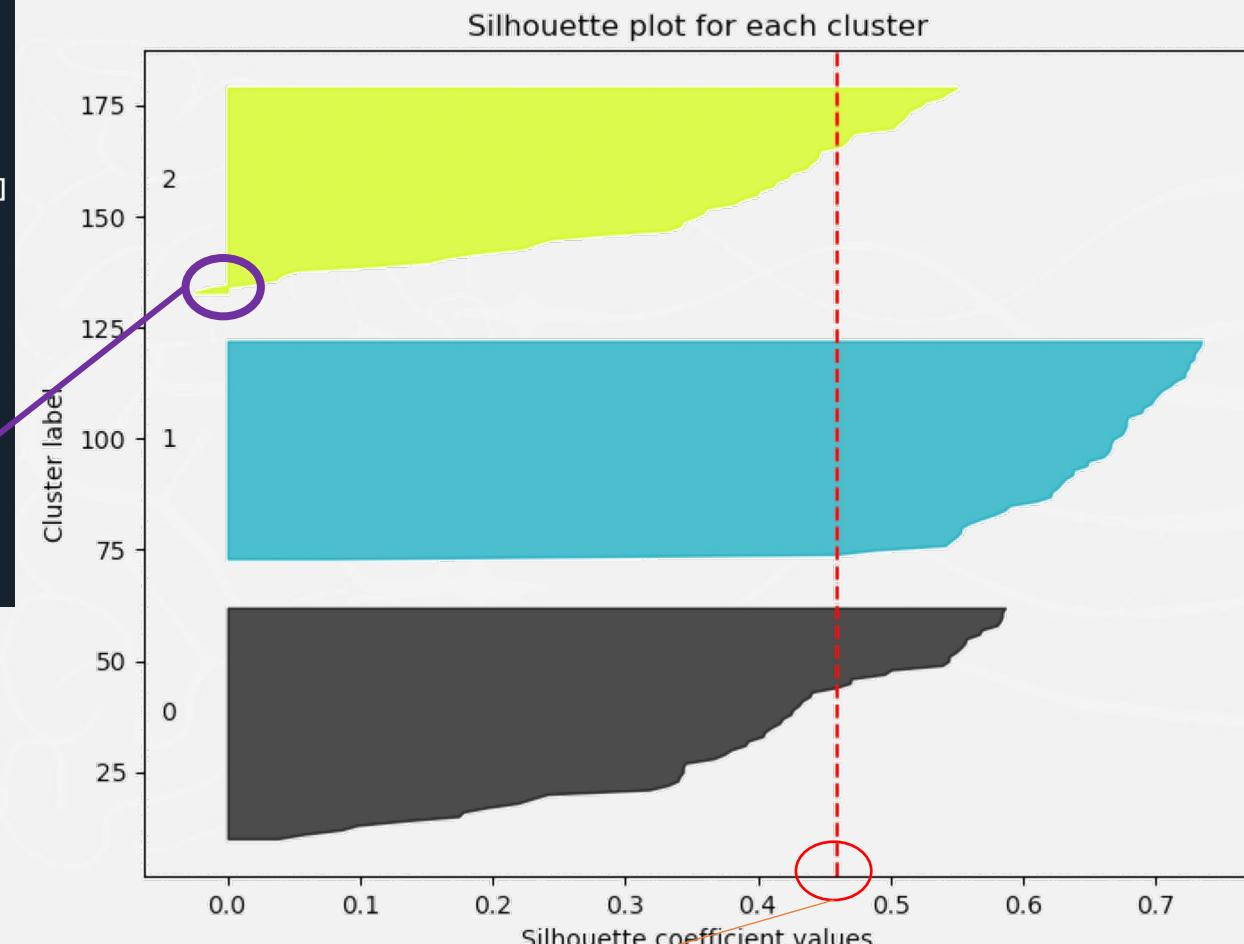
```
32 """
33 Get silhouette scores for each sample
34 """
35 silhouette_avg = silhouette_score(X_scaled, cluster_labels)
36 """
37 Plot the clustering regions and the centroids for all 2D combinations
38 """
39 #get all combinations of 2 features from the dataset
40 feature_combinations = list(combinations(range(X_scaled.shape[1]), 2))
41 #create subplots for all feature combinations
42 fig, axes = plt.subplots(len(feature_combinations) // 2, 2, figsize=(8, 8))
43 axes = axes.ravel()
44 for i, (feature1, feature2) in enumerate(feature_combinations):
45     ax = axes[i]
46     #plot the clustering regions in 2D for the current feature pair
47     colors = cm.nipy_spectral(cluster_labels.astype(float) / n_clusters)
48     ax.scatter(
49         X_scaled[:, feature1], X_scaled[:, feature2],
50         marker='o', s=100, c=colors, edgecolor='k'
51     )
52     #plot the centroids of each cluster
53     cluster_centers = kmeans.cluster_centers_
54     ax.scatter(
55         cluster_centers[:, feature1], cluster_centers[:, feature2],
56         marker='x', s=200, c='red', label='Centroids'
57     )
58     ax.set_xlabel(f'Feature {feature1 + 1}')
59     ax.set_ylabel(f'Feature {feature2 + 1}')
60     ax.legend()
61     #adjust and show the subplots
62     plt.tight_layout()
63     plt.show()
```



```

64 """
65 Plot the silhouette scores for each sample
66 """
67 plt.figure(figsize=(8, 6))
68 sample_silhouette_values = silhouette_samples(X_scaled, cluster_labels)
69 y_lower = 10
70 for i in range(n_clusters):
71     ith_cluster_silhouette_values = sample_silhouette_values[cluster_labels == i]
72     ith_cluster_silhouette_values.sort()
73     size_cluster_i = ith_cluster_silhouette_values.shape[0]
74     y_upper = y_lower + size_cluster_i
75     color = cm.nipy_spectral(float(i) / n_clusters)
76     plt.fill_betweenx(np.arange(y_lower, y_upper), 0,
77                       ith_cluster_silhouette_values, facecolor=color,
78                       edgecolor=color, alpha=0.7)
79     plt.text(-0.05, y_lower + 0.5 * size_cluster_i, str(i))
80     y_lower = y_upper + 10
81 plt.axvline(x=silhouette_avg, color="red", linestyle="--")
82 plt.title("Silhouette plot for each cluster")
83 plt.xlabel("Silhouette coefficient values")
84 plt.ylabel("Cluster label")
85 plt.show()

```

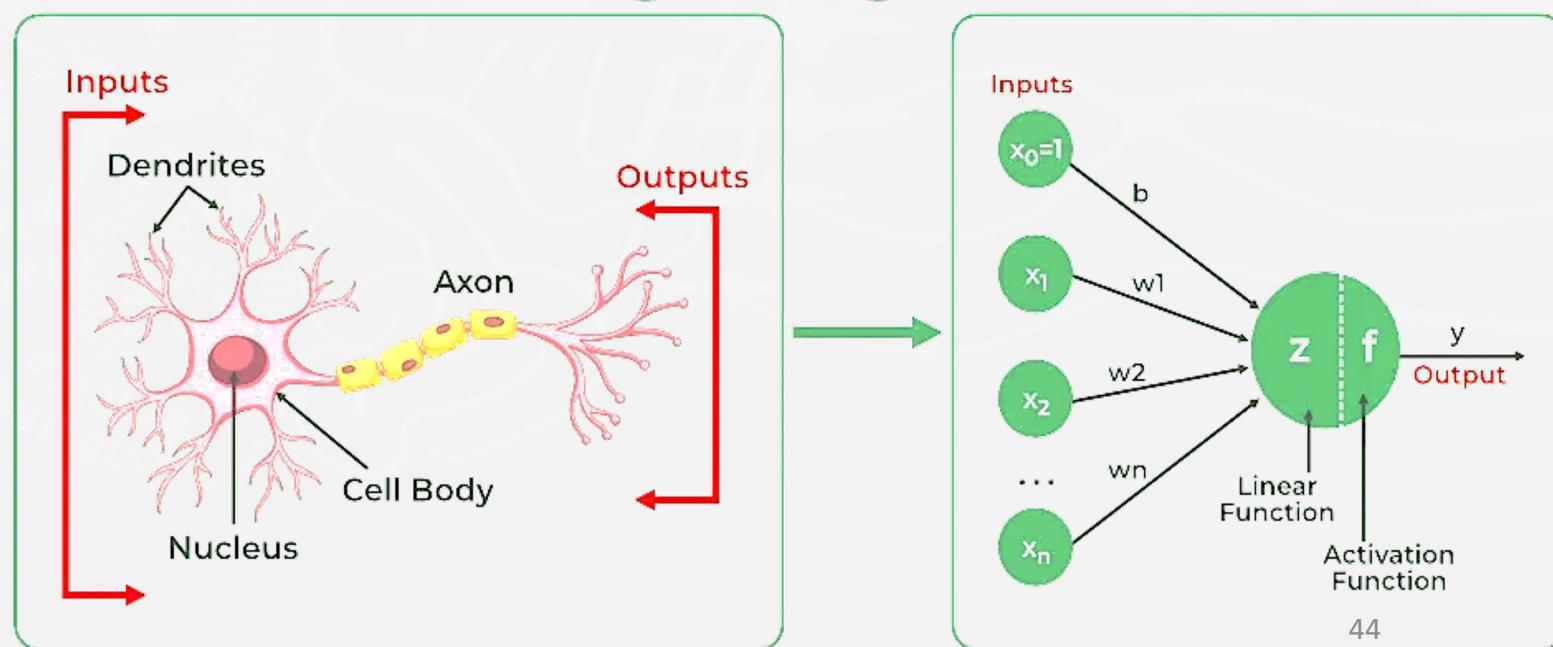
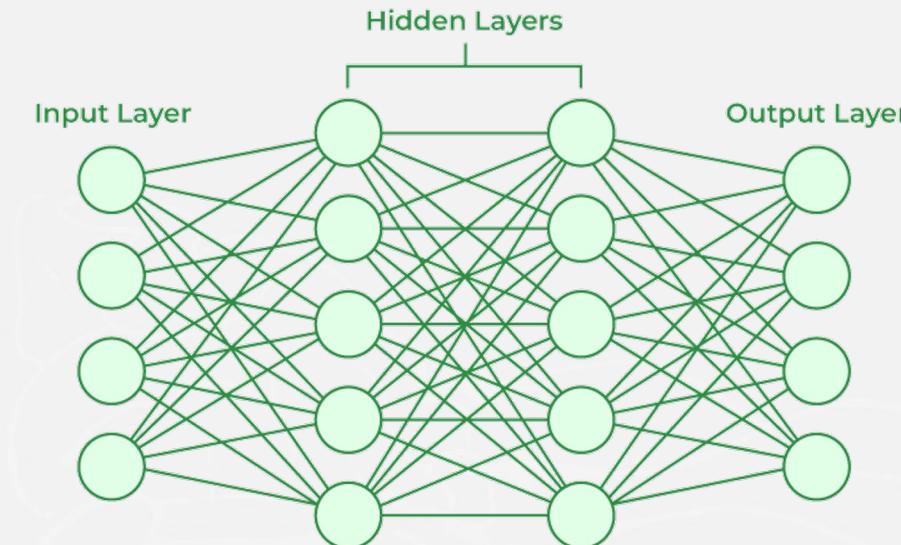


**silhouette\_avg**  
0.45994823920518635

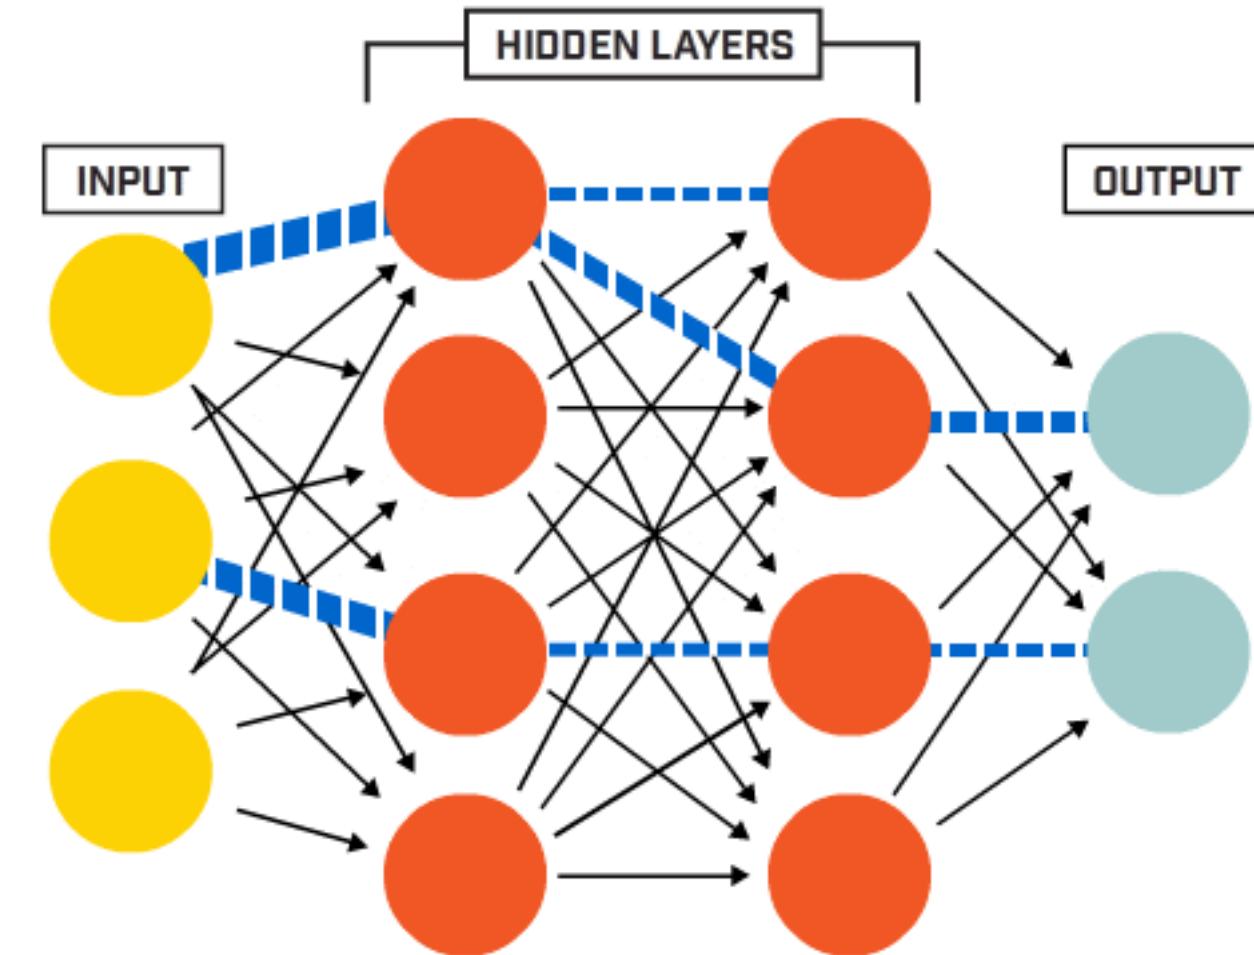
## Regression - Artificial neural network (ANN)

Comprised of a node layers,  
where each node connects to  
another and has an associated  
weight and threshold.

Optimize the weights by using  
corrective feedback.



## Regressions - ANN



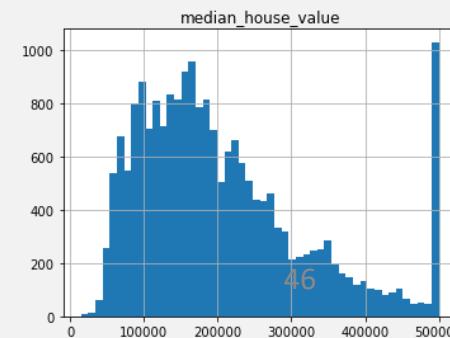
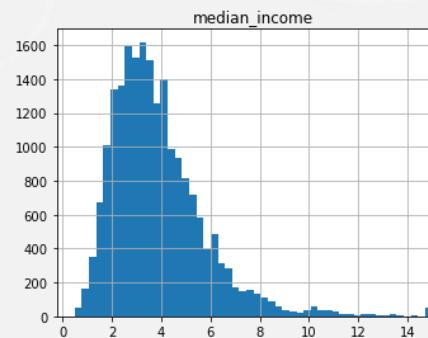
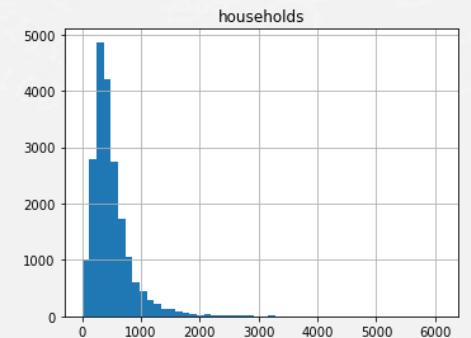
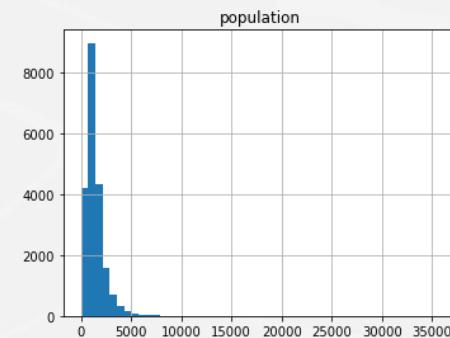
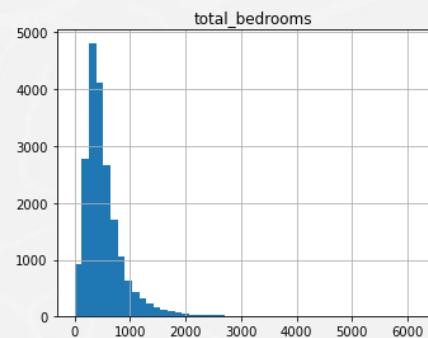
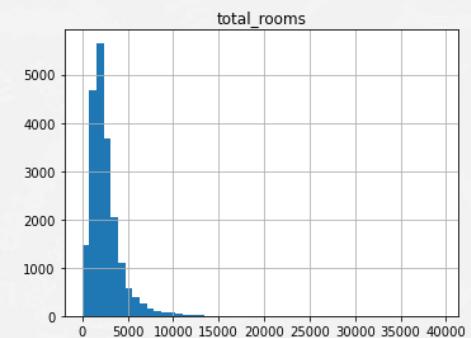
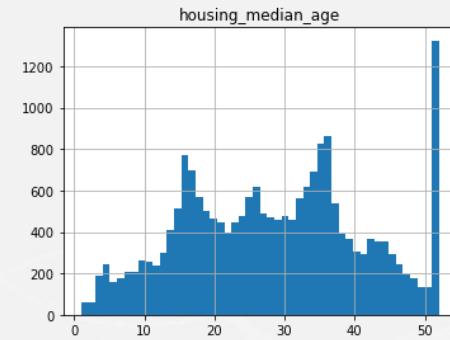
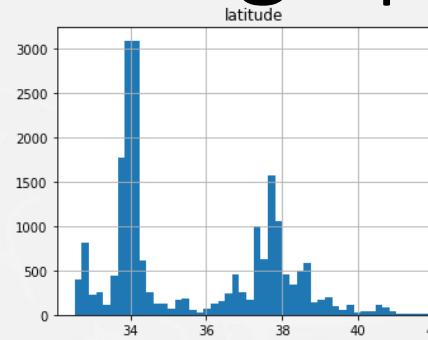
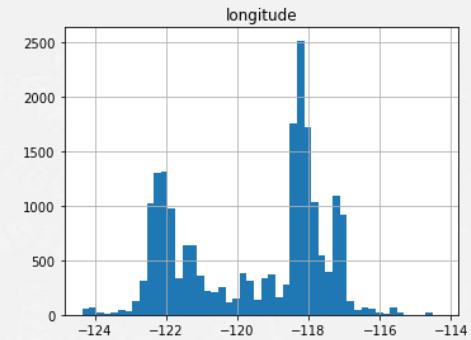
## 2- Examples of machine learning applications

### House price prediction California housing Dataset

Target variable – median housing value

Features:

- Median income in block
- Median house age in block
- Average number of rooms per house
- Average number of bedrooms per house
- Population of block
- Average number of household members per house
- Latitude of block
- Longitude of block



## 2- Examples of machine learning applications

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

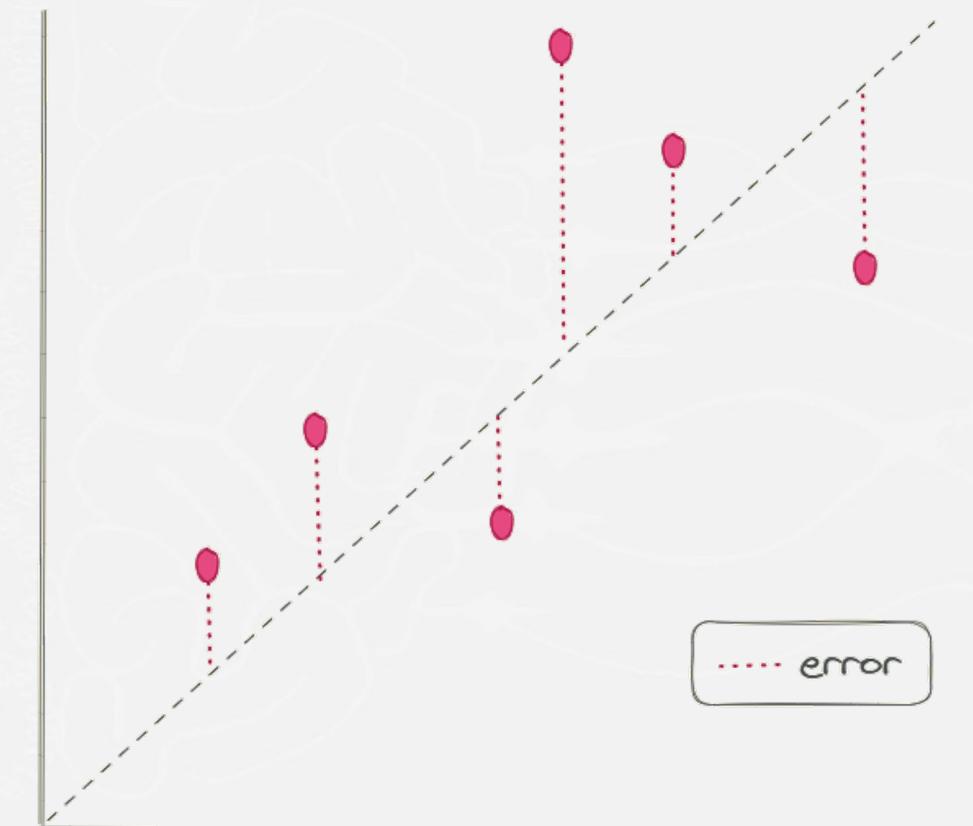
$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

$$R^2 = 1 - \frac{SSR}{SST} = 1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y})^2}$$

$$MAPE = \frac{1}{n} \sum_{i=1}^n \left| \frac{y_i - \hat{y}_i}{y_i} \right|$$

Predicted value



Actual value

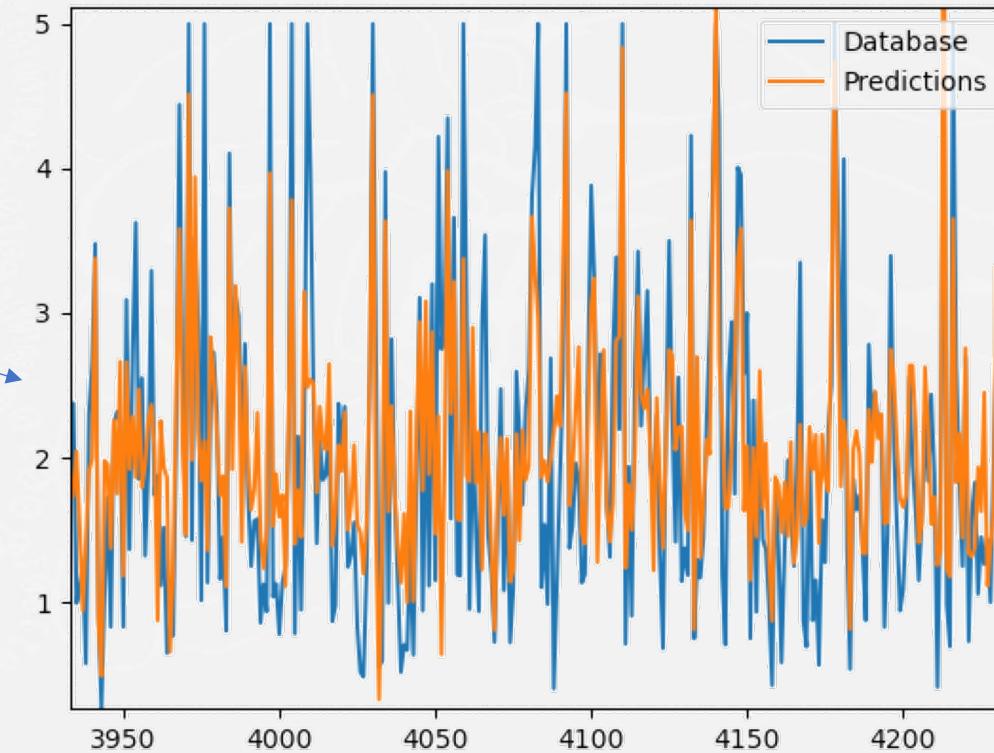
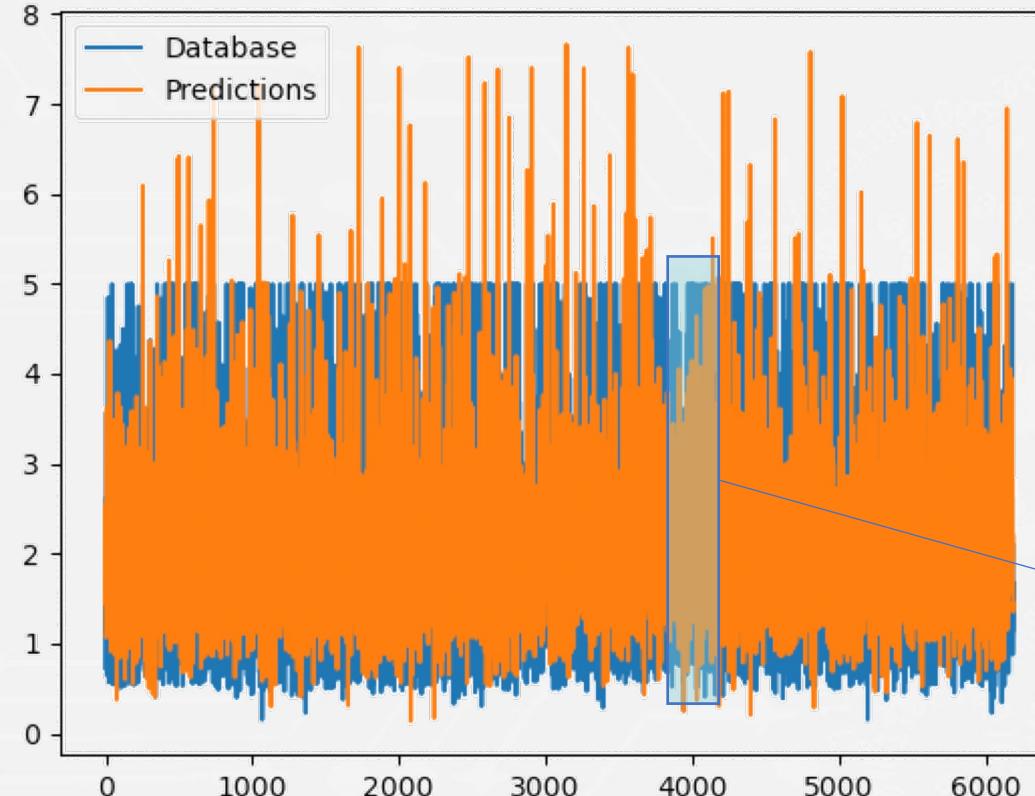
# Regression - ANN

```
1  """
2  Import libraries
3  """
4  from sklearn.neural_network import MLPRegressor
5  from sklearn import datasets
6  from sklearn.model_selection import train_test_split
7  from sklearn.datasets import fetch_california_housing
8  import numpy as np
9  from sklearn.metrics import mean_squared_error
10 import matplotlib.pyplot as plt
11 """
12 Import the dataset
13 """
14 housing = fetch_california_housing()
15 X = housing.data
16 y = housing.target
17 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
18 """
19 Define the model
20 """
21 nn = MLPRegressor(
22     activation='relu',
23     hidden_layer_sizes=(1000),
24     alpha=0.01,
25     early_stopping=True
26 )
27 """
28 Train the model
29 """
30 nn.fit(X_train, y_train)
```

```
31 """
32 Make predictions and estimate the error
33 """
34 pred = nn.predict(X_test)
35 test_set_rsquared = nn.score(X_test, y_test)
36 test_set_rmse = np.sqrt(mean_squared_error(y_test, pred))
37 print('Rsquared value: ', test_set_rsquared)
38 print('RMSE: ', test_set_rmse)
39 """
40 Plot the predictions
41 """
42 plt.plot(y_test, label = "Database")
43 plt.plot(pred, label = "Predictions")
44 plt.legend()
45 plt.show()
```

```
Rsquared value: 0.5144754879313564
RMSE: 0.7989435898734182
```

## Regression - ANN



## 3- Key remarks

Machine learning algorithms **learn from data** to find hidden relations and make predictions.

**A machine learning algorithm is as good as its input data.**

The algorithms require proper optimization and validation methods to ensure the generalization capability of the model.

